

# **Lab Lecture**

# **SQL Support for OLAP**

# SQL Extensions for Complex Data Analysis

Early versions of SQL provide limited support for complex data analysis

## Processing scenario

- retrieve data from the database into the client application environment
- perform analysis on the client
- overhead by retrieval of huge amounts of data

Data analysis features of SQL:99

## SQL/OLAP

- aggregate functions
- user-defined functions
- GROUP BY and HAVING
- CUBE and ROLLUP
- Window

# SQL:1992 Select – An example

**SELECT**

Country.region, Time.year, Manufacturer, SUM(salesAmt)

**FROM**

Sales, Time, Product, Country

**WHERE**

Sales.product\_ID = Product.product\_ID

AND Sales.time\_ID = Time.time\_ID

AND Sales.country\_ID = Country.country\_ID

AND Product.product\_category = 'Fridge'

AND Country.country\_name = 'New Zealand'

AND Time.year = 2011

**GROUP BY**

Country.region, Time.year, Manufacturer

**Relational treatment of multi-dimensional queries:**

- o depends on Schema - Star vs. Snowflake
- o joins between n-dimension-tables and the fact-table
- o restrictions are on the dimension-table

# Aggregation (Subtotals and Totals)

```
SELECT Product, NULL, NULL, SUM(turnover)
FROM Sales
WHERE Product = 'Fridge'
GROUP BY product
```

**UNION**

```
SELECT Product, Year, NULL, SUM(turnover)
From Sales
WHERE Product = 'Fridge'
GROUP BY Product, Year
```

**UNION**

```
SELECT Product, Year, Region, SUM(turnover)
FROM Sales
WHERE Product = 'Fridge'
GROUP BY Product, Year, Region
```

# Calculation of Subtotals and Totals

Product	Year	Region	Turnover per year and region	Turnover per year	Turnover
Fridge	2005	Northland	135		
		Auckland	120		
				255	
	2006	Northland	140		
		Auckland	135		
				275	
					530

Product	Year	Region	Turnover
Fridge	2005	Northland	135
Fridge	2005	Auckland	120
Fridge	2005		255
Fridge	2006	Northland	140
Fridge	2006	Auckland	135
Fridge	2005		275
Fridge			530

# Disadvantages of the UNION Version

Not efficient!

- many sub queries are necessary for calculation of subtotal
- grouping attributes specified more than once
- Join operations may have to be done more than once

Awkward to define!

- There might be an easier way
  - by using SQL:99 (SQL/OLAP) extensions

# The ROLLUP Operator

- ROLLUP is a modification ('enlargement') of the GROUPBY Clause
- ROLLUP produces subtotals
  - delivers subtotals for each aggregation level, up to the total
- ROLLUP greatly simplifies the amount of computation

**SELECT ... GROUP BY ROLLUP(columnList)**

# The ROLLUP Operator

```
SELECT Company, Year, SUM(quantity) AS Sum  
FROM Sales ...  
GROUP BY ROLLUP (Company, year)
```

Company	Year	Sum	
Siemens	1999	2.000	
Siemens	2000	3.000	
Siemens	2001	3.500	
Motorola	1999	1.000	
Motorola	2000	1.000	
Motorola	2001	1.500	
Nokia	1999	1.000	
Nokia	2000	1.500	
Nokia	2001	2.000	
Siemens	NULL	8.500	↑
Motorola	NULL	3.500	↓
Nokia	NULL	4.500	↑
NULL	NULL	16.500★	↓

**GROUP BY Company, Year**

**GROUP BY Company**

Without **GROUP BY Company**

- ROLLUP simplifies the amount of computation.
- Only one sort is needed.
- Thereafter, only aggregate (sum) operations are required.

# The CUBE Operator

- The Cube-Operator is a generalization of the **ROLLUP**operator
- It computes all combinations of aggregates across the dimensions included
- The sequence of operation is not dependent on the sequence given in the query

# CUBE Operator - Example

```
SELECT Producer, Year, SUM(quantity) AS Total  
FROM Sales  
Group BY Cube (Producer, Year)
```

The aggregation levels are:

1. Producer, Year
2. Producer
3. Year
4. Total

# CUBE Operator – Example...

**SELECT** Producer, Year, Sum (quantity) as Total

**FROM** Sales

**Group BY Cube** (Producer, Year)

The aggregation levels are:

1. Producer, Year
2. Producer
3. Year
4. Grand Total

Producer	Year	Total
Siemens	1999	2000
Siemens	2000	3000
Siemens	2001	3500
Motorola	1999	1000
Motorola	2000	1000
Motorola	2001	1500
Nokia	1999	1000
Nokia	2000	1500
Nokia	2001	2000
Siemens	NULL	8500
Motorola	NULL	3500
Nokia	NULL	4500
NULL	1999	4000
NULL	2000	5500
NULL	2001	7000
NULL	NULL	16500

**GROUP BY**  
(Producer,Year)

**GROUP BY**  
(Producer)

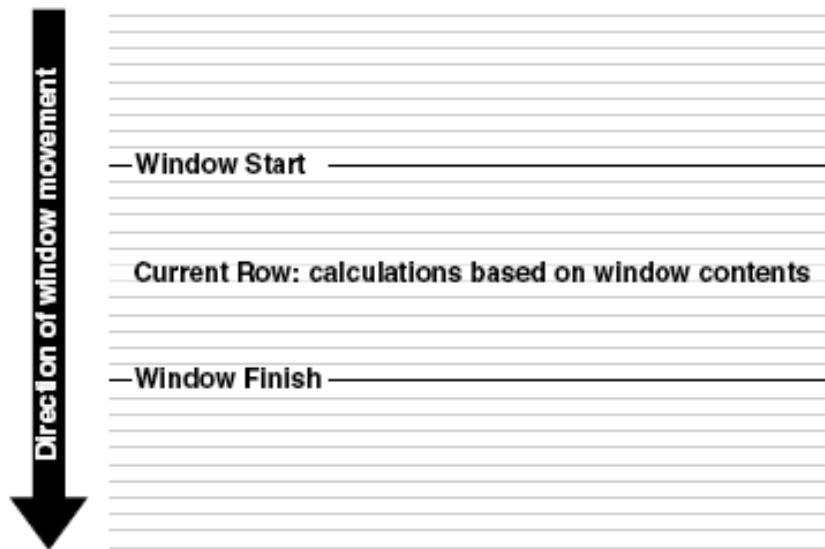
**GROUP BY**  
(Year)

Without  
**GROUP BY**  
clause

# The WINDOW clause

- Intuitively identifies an ordered ‘window’ of rows ‘around’ each row in a table.
  - Allows us to apply a rich collection of aggregate functions (SUM, AVG) to the window of a row and extend the row with the results

*Figure 21–2 Sliding Window Example*



Source: Oracle Data  
Warehousing Guide

# GROUP BY v WINDOW

- GROUPBY clause allows us to:
  - Create groupings (partitions) of rows
  - Apply aggregates to groups
  - However,
    - There is a single output row per partition
    - Rows within each partition are unordered
- A window is a combination of rows to which a function is applied, giving us:
  - cumulative totals
  - sliding averages
  - calculation over a partition or range

# Window example

```
SELECT L.State, T.month, AVG(S.Sales) OVER W AS movavg  
FROM Sales S, Time T, Locations L  
WHERE S.timeID = T.timeID  
AND S.LocationID = L.LocationID  
WINDOW W AS (PARTITION BY L.State  
              ORDER BY T.Month  
              RANGE BETWEEN INTERVAL '1' MONTH  
                      PRECEDING  
              AND INTERVAL '1' MONTH FOLLOWING)
```

# 3 Steps in Defining a Window

1. Define *partitions* of the table
  - Partitions are based on the L.state column
2. Specify the *ordering* of rows within a partition
  - Rows within each partition are ordered by T.month
3. Frame windows – establish the boundaries of the window associated with each row in terms of the ordering of rows within partitions
  - Window for a row includes the row itself plus
  - all rows whose month value is within a month before or after
  - ie. A row whose month value is June 2009 has a window containing all rows with *month* equal to May, June or July 2009
  - The answer row corresponding to a given row is constructed by first identifying its window
  - Then, for each answer column (defined using a window aggregate function), the aggregate is computed using the rows in the window

# Framing a Window

Two ways to frame a window in SQL:99

1. Using a RANGE construct
2. Directly specifying how many rows before and after the given row are in its window

```
WINDOW W AS(PARTITION BY L.State  
            ORDER BY T.Month  
            RANGE BETWEEN 1 PRECEDING  
                  AND 1 FOLLOWING)
```

# Framing a Window

- A window can consist of **all rows of a partition** or it can be defined as a “**gliding window**” with the smallest size of **one row** of the partition.
- The window defines the selection of rows for calculation.
  - The actual row serves as reference point for the starting and ending point of a window.
- The size of a window can be defined physically (**ROWS**) by defining the number of rows before and after the actual row.
- The size of a window can be defined logically by values in relation to the actual row (see **RANGE**definition)
- A window has a starting and an ending row.
- Depending on the definition a window can move on both sides. Such a “**moving window**” is defined by **UNBOUNDED PRECEDING** and **UNBOUNDED FOLLOWING** e.g.:
  - A window for a cumulative sum **has a starting row**, and can glide from there to the last row of the partition.
  - A window for a gliding average **has a starting and an ending row**, and can glide over the whole partition

# Sliding Window with a gliding average

```
SELECT Producer, Year, Sum(Sales) AS Total, Avg(Total) OVER w AS Moving_Avg  
FROM SalesFact  
WHERE Year BETWEEN 1999 and 2001  
WINDOW w AS(PARTITION BY Producer  
            ORDER BY Year  
            ROWS 1 PRECEDING)
```

Producer	Year	Total	Moving_Avg
Motorola	1999	1000	1000
Motorola	2000	1000	1000
Motorola	2001	1500	1250
Nokia	1999	1000	1000
Nokia	2000	1500	1250
Nokia	2001	2000	1750
Siemens	1999	2000	2000
Siemens	2000	3000	2500
Siemens	2001	3500	3250

# Window

```
SELECT country, month, group, sales,  
       SUM(sales) OVER w AS moving_sum  
FROM Sales AS s  
WINDOW w AS (
```

PARTITION BY s.country, s.group  
ORDER BY s.month ← window ordering  
ROWS 2 PRECEDING ) ← window framing

```
SELECT country, month, group, sales,  
       SUM(sales) OVER (  
           PARTITION BY s.country,  
                     s.group  
           ORDER BY s.month  
           ROWS 2 PRECEDING )  
           AS moving_sum  
FROM Sales AS s
```

window partitioning

window ordering

window framing



in-line window  
specification

Calculated column

country	month	group	sales	moving_sum
F	01/04	100	500	500
F	02/04	100	1000	1500
F	03/04	100	250	1750
F	04/04	100	750	2000
F	02/04	200	1250	1250
F	03/04	200	2000	3250
F	04/04	200	500	3750
GB	01/04	100	400	400
GB	02/04	100	800	1200
GB	03/04	100	300	1500
GB	04/04	100	2000	3100
GB	05/04	100	2500	4800
GB	02/04	200	100	100
GB	03/04	200	100	200
GB	04/04	200	100	300
GB	05/04	200	100	300
I	01/04	300	250	250
I	02/04	300	750	1000
I	03/04	300	200	1200
I	04/04	300	1500	2450
I	05/04	300	800	2500

# Window Framing

```
...ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING
/* physical aggregation group on a dense column */

... ROWS BETWEEN 5 PRECEDING AND CURRENT ROW ... ROWS 5 PRECEDING
/* preceding 5 rows and current row */ ... ROWS UNBOUNDED PRECEDING
/* all preceding rows and current row */

... ROWS BETWEEN CURRENT ROW AND UNBOUNDED FOLLOWING
/* current row and all following rows */

... ROWS BETWEEN 6 PRECEDING AND 3 PRECEDING
/*the 3 rows starting 6 rows before the current row, current row is not included */
```

```
... ORDER BY month RANGE BETWEEN INTERVAL '1' MONTH PRECEDING
AND INTERVAL '1' MONTH FOLLOWING
/* logical aggregation group including the last, current, and next month */

... ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING
/* physical aggregation group on a dense column
```

# Ranking Functions

- Allows you to rank items in a group (window, partition)
  - Eg. Finding the top 3 products sold in Auckland last year
- Three functions
  - **RANK**
  - **DENSE\_RANK**
  - **PERCENT\_RANK**
- Example: Ranking a competition where 3 people have tied for 2nd place
  - Using **RANK**  
**1, 2, 2, 2, 5,...**
  - Using **DENSE\_RANK** leaves no gaps  
**1, 2, 2, 2, 3,...**

# RANK OPERATORS

- Example

```
SELECT Producer, Year, Sum(Sales),  
RANK() OVER(ORDER BY Sum(Sales) DESC) AS Rank  
FROM SalesFact  
WHERE Year = 2010  
GROUP BY Producer
```

Producer	Year	Sum(Sales)	Rank
Siemens	2001	3500	1
Nokia	2000	2000	2
Motorola	1999	1500	3

- To produce the table on the next slide:

```
SELECT Producer, Year, Sum(Sales),  
RANK() OVER(ORDER BY Sum(Sales) DESC) AS Rank DENSE_RANK()  
OVER(ORDER BY Sum(Sales) DESC) AS Drank PERCENT_RANK()  
OVER(ORDER BY Sum(Sales) DESC) AS PCT_Rank  
FROM SalesFact  
WHERE Year BETWEEN 2008 and 2010  
GROUP BY Producer
```

# RANK OPERATORS

DENSE\_RANK calculates a rank as well but DOES NOT produce gaps

Producer	Year	Total	Rank	Drank	Pct_Rank
Siemens	2010	3500	1	1	0
Siemens	2009	3000	2	2	0.17
Nokia	2010	2000	3	3	0.33
Siemens	2008	2000	3	3	0.33
Nokia	2009	1500	5	4	0.67
Motorola	2010	1500	5	4	0.67
Nokia	2008	1000	7	5	1.00
Motorola	2009	1000	7	5	1.00
Motorola	2008	1000	7	5	1.00

# References

- Ramakrishnan & Gherke “Database Management Systems”, 3<sup>rd</sup> edition.
- Oracle Data Warehousing Guide 10g, Chapters 20, 21.