
Lecture 08

Optimizations

Summary – last week

Summary

- B-Trees are not fit for multidimensional data
- R-Trees
 - MBR as geometry to build multidimensional indexes
 - Inefficient because they allow overlapping between neighboring MBRs
 - R⁺-trees - improve the search performance
- UB-Trees
 - Reduce multidimensional data to one dimension in order to use B-Tree indexes
 - Z-Regions,Z-Curve,use the advantage of bit operations to make optimal jumps
- Bitmap indexes
 - Great for indexing tables with set-like attributes e.g., Gender:Male/Female
 - Operations are efficient and easy to implement (directly supported by hardware)



This week

Optimization

1. Partitioning
2. Joins
3. MaterializedViews



Partitioning

- Breaking the data into several **physical units** that can be handled separately
- Granularity and partitioning are key to **efficient implementation of a warehouse**
- The question is not whether to use partitioning but how to do it



Partitioning (cont'd.)

- Why partitioning?
 - **Flexibility** in managing data
 - Smaller physical units allow
 - Easy restructuring
 - Free indexing
 - Sequential scans, if needed
 - Easy reorganization
 - Easy recovery
 - Easy monitoring



Partitioning (cont'd.)

- In DWs, partitioning is done to improve:
 - **Business query performance**, i.e., minimize the amount of data to scan
 - **Data availability**, e.g., back-up/restores can run at the partition level
 - **Database administration**, e.g., adding new columns to a table, archiving data, recreating indexes, loading tables

Partitioning (cont'd.)

- Possible approaches:
 - **Data partitioning**
where data is usually partitioned by
 - Date
 - Line of business
 - Geography
 - Organizational unit
 - Combinations of these factors
 - **Hardware partitioning**
 - Hardware partitioning makes data available to different processing nodes by ensuring that sub-processes are capable of running on the different nodes



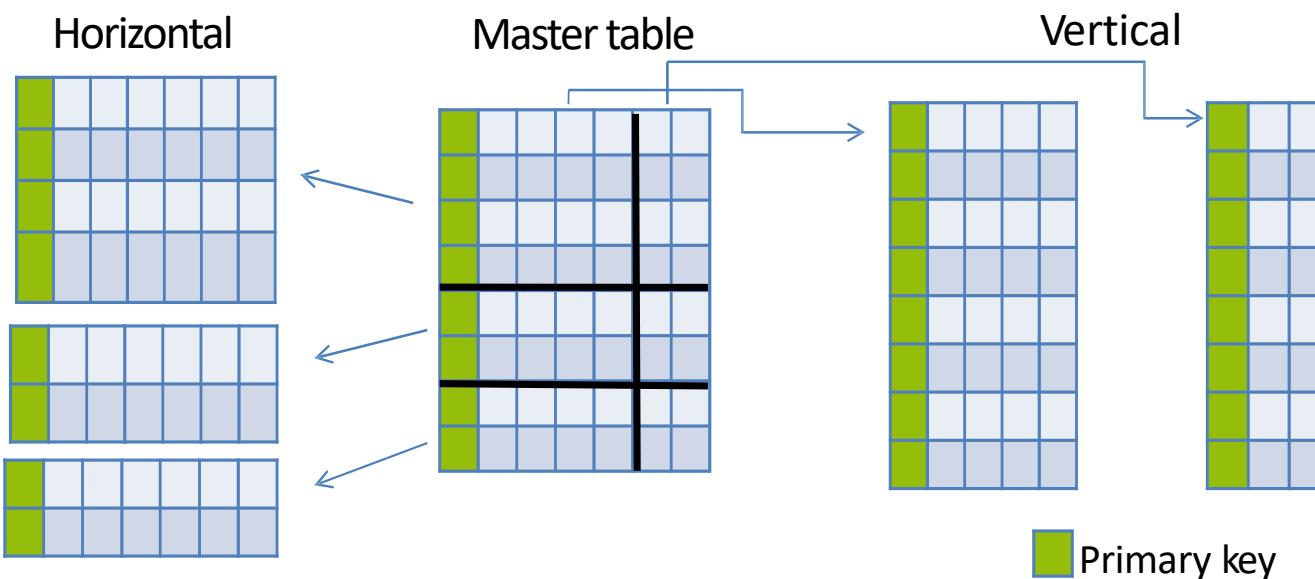
Partitioning (cont'd.)

- At which **level** should data partitioning occur?
 - Possibilities are
 - on application
 - or DBMS level
- Partitioning on DBMS level is clear; but it also makes sense to partition at **application level**
 - E.g., allows different definitions for each year
 - Important since DWs spans many years and as business evolves DWs change, too
 - Think for instance about changing tax laws



Data Partitioning

- Data partitioning, involves:
 - Splitting out the rows of a table into multiple tables i.e., **horizontal partitioning**
 - Splitting out the columns of a table into multiple tables i.e., **vertical partitioning**



Which partitioning type is good tax laws changing every year?

Data Partitioning (cont'd.)

- Horizontal partitioning
 - The set of tuples of a table is split among **disjoint** table parts
 - Definition: A set of Relations $\{R_1, \dots, R_n\}$ represent a **horizontal partitioning** of Master-Relation R if and only if $R_i \subseteq R$, $R_i \cap R_j = \emptyset$ and $R = \bigcup_i R_i$, for $1 \leq i, j \leq n$
 - According to the partitioning procedure we have different horizontal partitioning solutions
 - Range partitioning, List partitioning and Hash partitioning

Horizontal Partitioning

- Range Partitioning
 - Selects a partition by determining if the partitioning key is inside a certain **range**
 - A partition can be represented as a restriction on the master-relation
 - $R_i = \sigma_{P_i}(R)$, where P_i is the partitioning predicate. The partitioning predicate can involve more attributes
 - $P_1: \text{Country} = \text{'Germany'}$ and $\text{Year} = 2009$
 - $P_2: \text{Country} = \text{'Germany'}$ and $\text{Year} < 2009$
 - $P_3: \text{Country} \neq \text{'Germany'}$

Horizontal Partitioning (cont'd.)

- List Partitioning
 - A partition is assigned a **list of values**. If the partitioning key has one of these values, the partition is chosen
 - For example all rows where the column Country is either Iceland, Norway, Sweden, Finland or Denmark could build a partition for the Scandinavian countries
 - Is also expressed as a simple **restriction** on the master relation
 - The partitioning predicate involves just one attribute
 - $P_1: \text{City IN } ('Hamburg', 'Hannover', 'Berlin')$
 - $P_2: \text{City IN } (\text{DEFAULT})$ – represents tuples which do not fit to P_1

Horizontal Partitioning (cont'd.)

- Hash Partitioning
 - The value of a hash function determines membership in a partition. Assuming there are four partitions, the hash function could return a value from 0 to 3
 - For each tuple t_j of the master-table R , the hash function will associate it to a partition table R_i
 - $R_i = \{t_1, \dots, t_m / t_j \in R \text{ and } H(t_j) = H(t_k) \text{ for } 1 \leq j, k \leq m\}$
 - This kind of partitioning is particularly used in parallel processing
 - The goal is to achieve an equal distribution of the data

Horizontal Partitioning (cont'd.)

- Horizontal partitioning in Data Warehousing partitions data by
 - Time dimension
 - Periods, such as week or month can be used or the data can be partitioned by the age of the data
 - E.g., if the analysis is usually done on last month's data the table could be partitioned into monthly segments
 - A dimension other than time
 - If queries usually run on a grouping of data: e.g., each branch tends to query on its own data and the dimension structure is not likely to change then partition the table on this dimension
 - On table size
 - If a dimension cannot be used, partition the table by a predefined size. If this method is used, metadata must be created to identify what is contained in each partition

Vertical Partitioning

- Vertical Partitioning
 - Involves creating tables with **fewer columns** and using additional tables to store the remaining columns
 - Different physical storage might be used e.g., storing infrequently used or very wide columns on a different device
 - Usually called **row splitting**
 - Row splitting creates a one-to-one relationship between the partitions

Vertical Partitioning (cont'd.)

- In DW, common vertical partitioning means
 - Moving **seldom used** columns from a highly-used table to another table
 - Creating a **view** across the two newly created tables restores the original table with a performance penalty
 - However, performance will increase when accessing the highly-used data e.g., for statistical analysis

Vertical Partitioning (cont'd.)

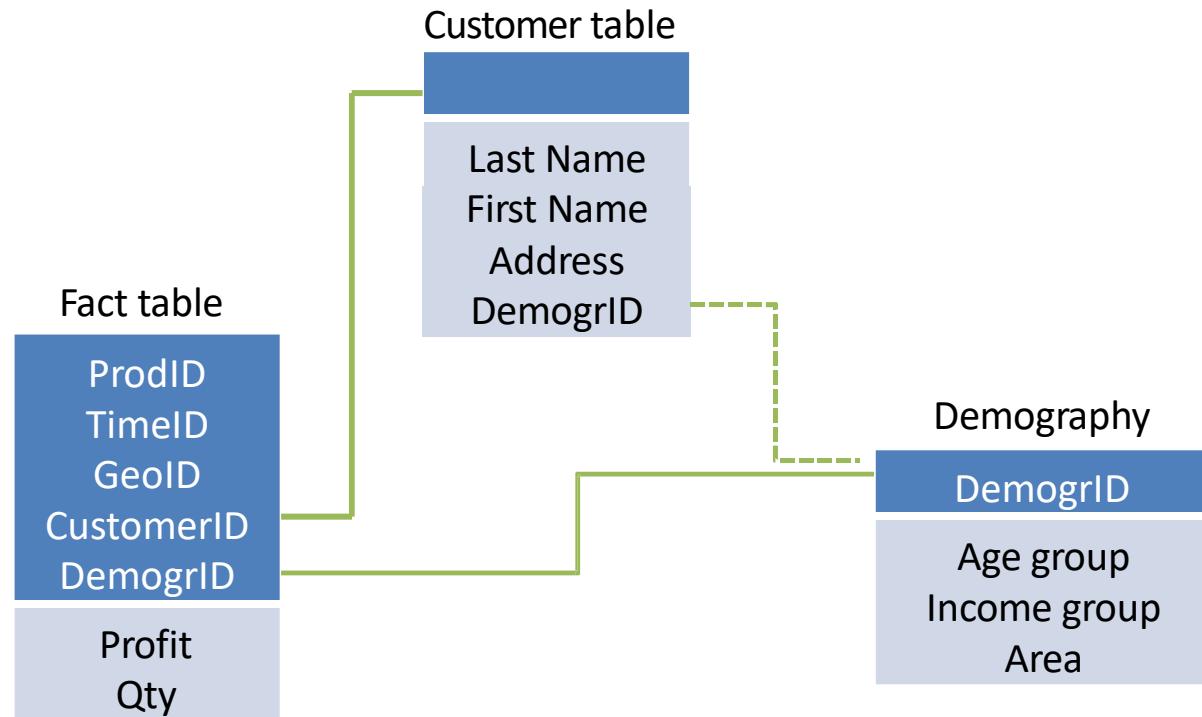
- In DWs with **very large** dimension tables (e.g., Amazon - a customer table with tens of millions of records) we have
 - Most of the attributes are rarely, if at all, queried
 - E.g., the address attribute is not as interesting as evaluating customers per age-group
 - But we must still maintain the link between the fact table and the **complete** customer dimension, which has **high performance costs!**

Vertical Partitioning (cont'd.)

- The solution is to use **Mini-Dimensions**, a special case of vertical partitioning
 - Many dimension attributes are used **very frequently** as browsing constraints
 - In big dimensions these constraints can be hard to find among the lesser used ones
 - Logical groups of often used constraints can be separated into **small dimensions** which are very well indexed and easily accessible for browsing

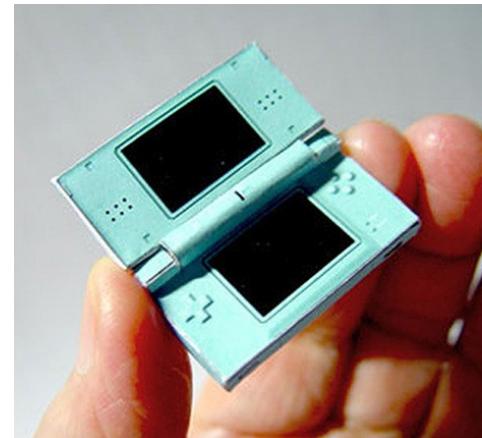
Vertical Partitioning (cont'd.)

- Mini-Dimensions, e.g., the **Demography** table



Vertical Partitioning (cont'd.)

- All variables in these mini-dimensions must be presented as **distinct classes**
- The key to the mini-dimension can be placed as a **foreign key** in both the fact table and dimension table from which it has been broken off
- Mini-dimensions, as their name suggests, should be kept small and compact

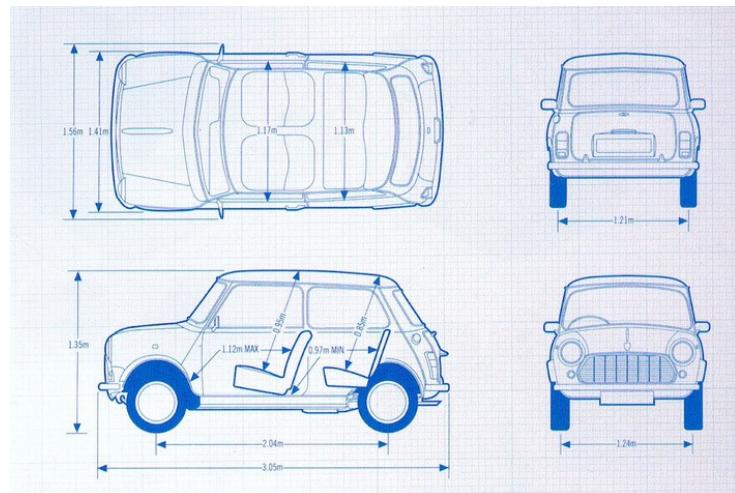


Partitioning

- Advantages
 - Records used together are grouped together
 - Each partition can be optimized for performance
 - Security, recovery
 - Partitions stored on different disks: contention
 - Take advantage of parallel processing capability
- Disadvantages
 - Slow retrieval across partitions (expensive joins)
 - Complexity

Partitioning (cont'd.)

- Use partitioning when:
 - A table is > 2GB (from Oracle)
 - A Table is > 100 Million rows (praxis)
 - Think about it, if table is > 1 million rows
 - Partitioning does not come for free!



Partitioning (cont'd.)

- **Partitioning management**
 - Partitioning should be transparent outside the DBMS
 - The applications work with the Master-Table at logical level
 - The conversion to the physical partition tables is performed internally by the DBMS
 - It considers also data consistency as if the data were stored in just one table
 - Partitioning transparency is not yet a standard. Not all DBMS support it!

Partitioning Management

- Partitions in practice

- Oracle supports Range-, List-, Hash-, Interval-, System-Partitions as well as combinations of these methods

- E.g., partitioning in Oracle:

- CREATETABLE SALES(

- ProdID NUMBER,

- GeolD NUMBER,

- TimeID DATE,

- Profit NUMBER)

- PARTITION BY RANGE(timeID)(

- PARTITION before 2008

- VALUES LESSTHAN (TO_DATE ('01-JAN-2008','DD-MM-YYYY')),

- PARTITION 2008

- VALUES LESSTHAN (TO_DATE ('01-JAN-2009','DD-MM-YYYY'))

-);



Partitioning Management (cont'd.)

- Partitions in practice
 - In Oracle partitioning is performed with the help of the LESSTHAN function. How can we partition data in the current year?
 - ALTER TABLE Sales

```
ALTER TABLE Sales ADD PARTITION after2008 VALUES LESS THAN  
(MAXVALUE);
```



Partitioning Management (cont'd.)

RowID	ProdID	GeOID	TimeID	Profit
...
121	132	2	05.2007	8K
122	12	2	08.2008	7K
123	15	1	09.2007	5K
124	14	3	01.2009	3K
125	143	2	03.2009	1,5K
126	99	3	05.2007	1K

RowID	ProdID	GeOID	TimeID	Profit
...
121	132	2	05.2007	8K
123	15	1	09.2007	5K
126	99	3	05.2007	1K

RowID	ProdID	GeOID	TimeID	Profit
122	12	2	08.2008	7K

RowID	ProdID	GeOID	TimeID	Profit
124	14	3	01.2009	3K
125	143	2	03.2009	1,5K

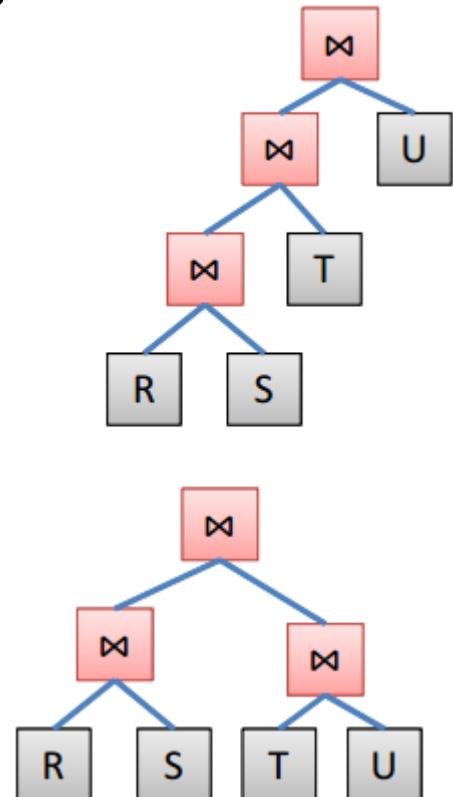
Join Optimization

- Often queries over **several** partitions are needed
 - This results in **joins** over the data
 - Though joins are **generally expensive operations**, the overall cost of the query may strongly differ with the chosen evaluation plan for the joins
- Joins are **commutative** and **associative**
 - $R \bowtie S \equiv S \bowtie R$
 - $R \bowtie (S \bowtie T) \equiv (S \bowtie R) \bowtie T$



Join Optimization (cont'd.)

- This allows to evaluate individual joins in **any order**
 - Results in **join trees**
 - Different join trees may show very different evaluation performance
 - Join trees have different **shapes**
 - Within a shape, there are different relation **assignments** possible
- Example: $R \bowtie S \bowtie T \bowtie U$



Join Optimization (cont'd.)

- Number of **possible join trees** grows rapidly with number of join relations
 - For n relations, there are $T(n)$ different tree shapes
 - $T(1) = 1$
 - $T(n) = \sum_{i=1}^{n-1} T(i)T(n-i)$
 - “Any number of $1 \leq i \leq n-1$ relations may be in the left subtree and ordered in $T(i)$ shapes while the remaining $n-i$ relations form the right subtree and can be arranged in $T(n-i)$ shapes.”

Join Optimization (cont'd.)

- Optimizer has 3 choices
 - Consider all possible join trees
 - Usually not possible
 - Consider a subset of all trees
 - i.e. restrict to trees of certain shapes
 - Use heuristics to pick a certain shape

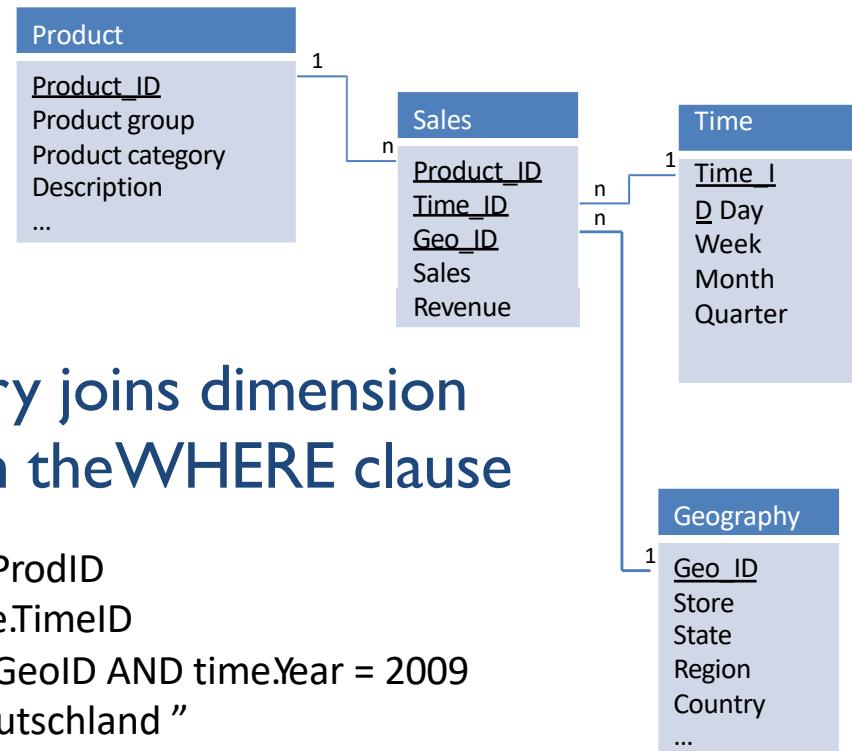
Join Optimization in DW

- Relational optimization of **star-joins**

- Star schema comprises a **big** fact table and many **small** dimension tables

- An **OLAP SQL** query joins dimension and fact tables usually in the WHERE clause

```
sales.ProdID = product.ProdID  
AND sales.TimeID = time.TimeID  
AND sales.GeoID = geo.GeoID AND time.Year = 2009  
AND geo.Country = "Deutschland"  
AND product.group = "wash machines"
```

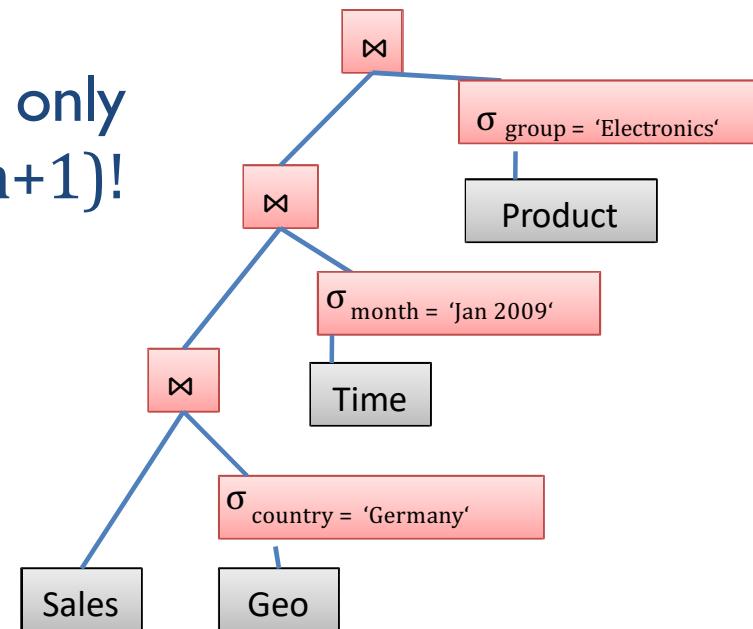


Join Optimization in DW (cont'd.)

- If the OLAP query specifies **restrictions** or **group by's** on n dimensions, an $n+1$ order join is necessary
 - Joins can be performed only pair-wise, resulting in $(n+1)!$ possible joining orders

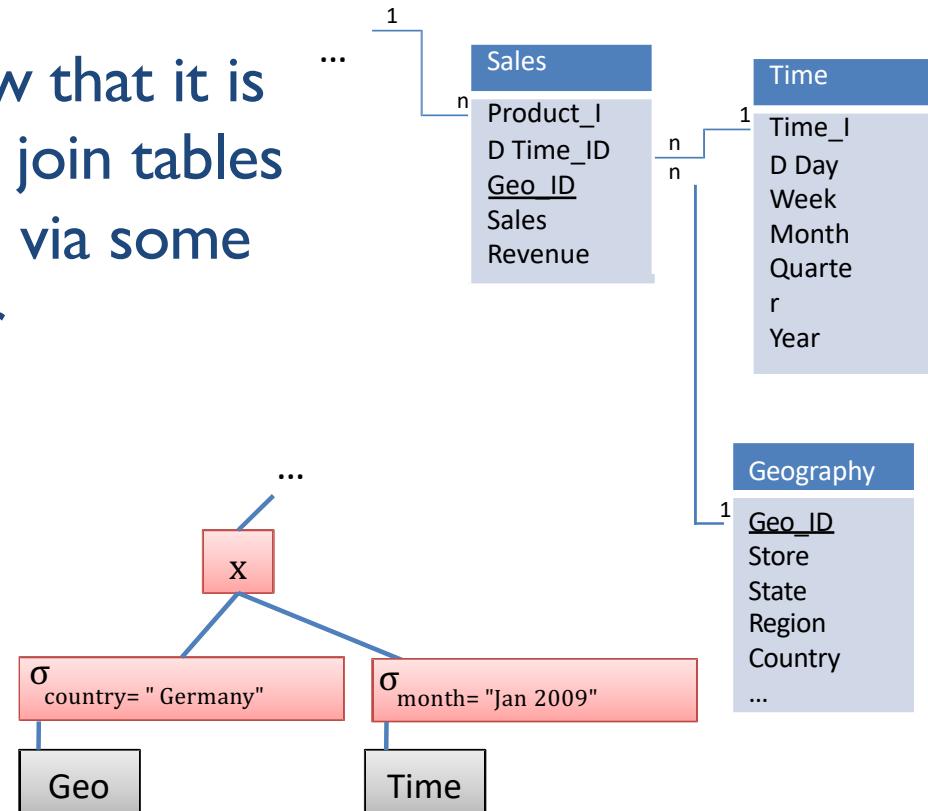
1. $\sigma_{country=Germany} (Sales \bowtie Geo)$

2. $Sales \bowtie (\sigma_{country=Germany} Geo)$



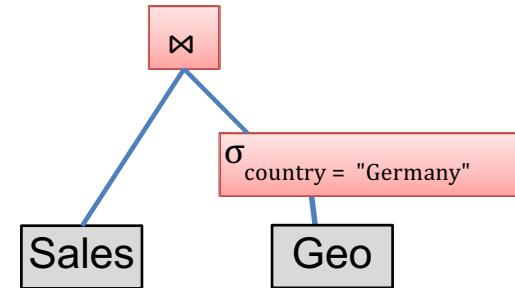
Join Heuristics

- To reduce the number of join-orders, **heuristics** are used
 - In OLTP heuristics show that it is **not a good idea** to join tables that are not connected via some attribute to each other
 - E.g., Geo with Time relation leads to a **Cartesian product**



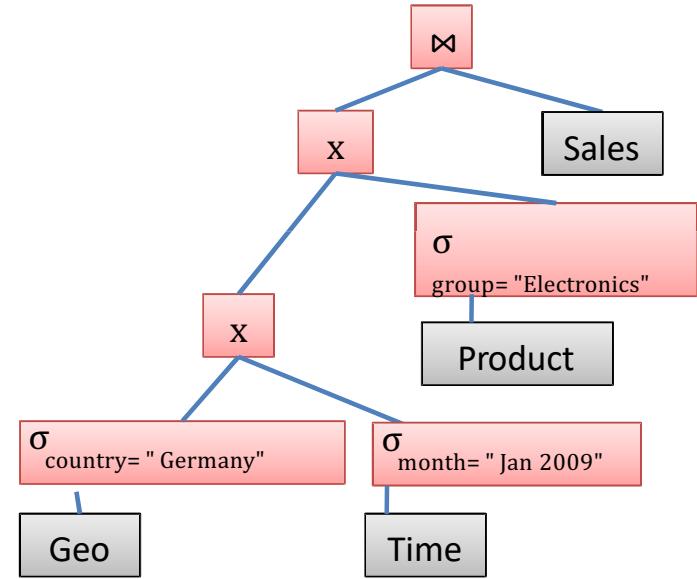
Join Heuristics (cont'd.)

- But this heuristic rule from OLTP is **not suitable** for DW!
 - E.g., join Sales with Geo in the following case:
 - Sales has 10 mil records, in Germany there are 10 stores, in January 2009 there were products sold in 20 days, and the Electronics group has 50 products
 - If 20% of our sales were performed in Germany, still we get 2mil intermediate results.
 - Also as we are interested in only one country from Geo dimension table so index on Geo table would not help that much either.



Dimensional Cross Product

- In star-joins a **cross product** of the dimension tables is recommended
 - Geo dimension – 10 stores
 - Time dimension 20 days
 - Product dimension 50 products
 - $10 \times 20 \times 50 = 10\,000$ records after performing the cross product of the dimensions
 - The selectivity is in this case 0.1% which is fit for using an index

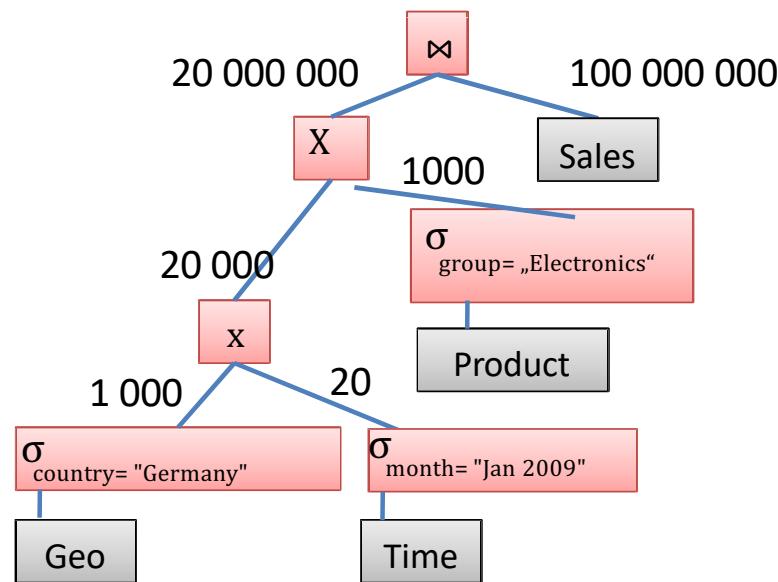


Dimensional Cross Product (cont'd.)

- But dimension cross products can also become **expensive**
 - If the restrictions on the dimensions are not restrictive enough or if there are many dimension tables
 - E.g., if we would have query on 1000 stores and 1000 customers, the cross product would comprise $1000 * 20 * 50 * 1000 = 100 \text{ mil records...}$

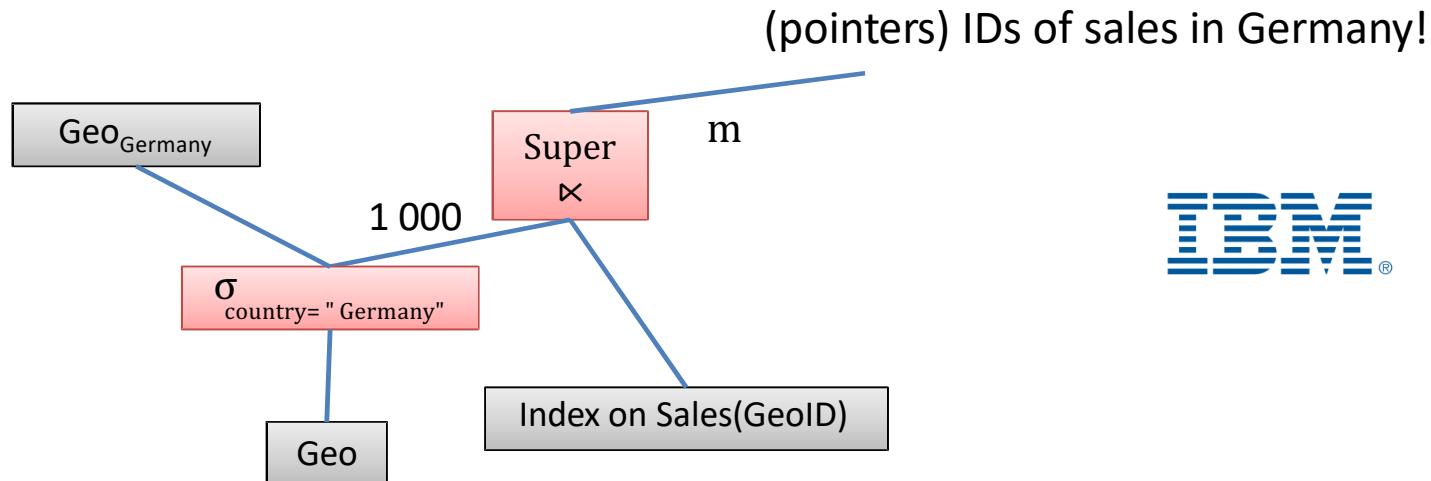
Dimensional Cross Product (cont'd.)

- E.g., In Germany the company has 1000 stores, it sold in 20 days in January, and it has 1000 types of electronics products
 - $1000 * 20 * 1000 = 20 \text{ mil records...}$



Star-join Optimization

- The **IBM DB2** solution for expensive dimension cross products is to build **semi-joins** of the dimension tables with indexes
 - A B*-Tree index will be on the fact table (sales) for each dimension (B*-Trees are unidimensional)



IBM®

Star-join Optimization (cont'd.)

Geo

Geoid	Store	City	Country
1	S1	BS	Germany
2	S2	BS	Germany
3	S3	HAN	Germany
4	S4	Lyon	France

Geoid	Store	City	Country
1	S1	BS	Germany
2	S2	BS	Germany
3	S3	HAN	Germany

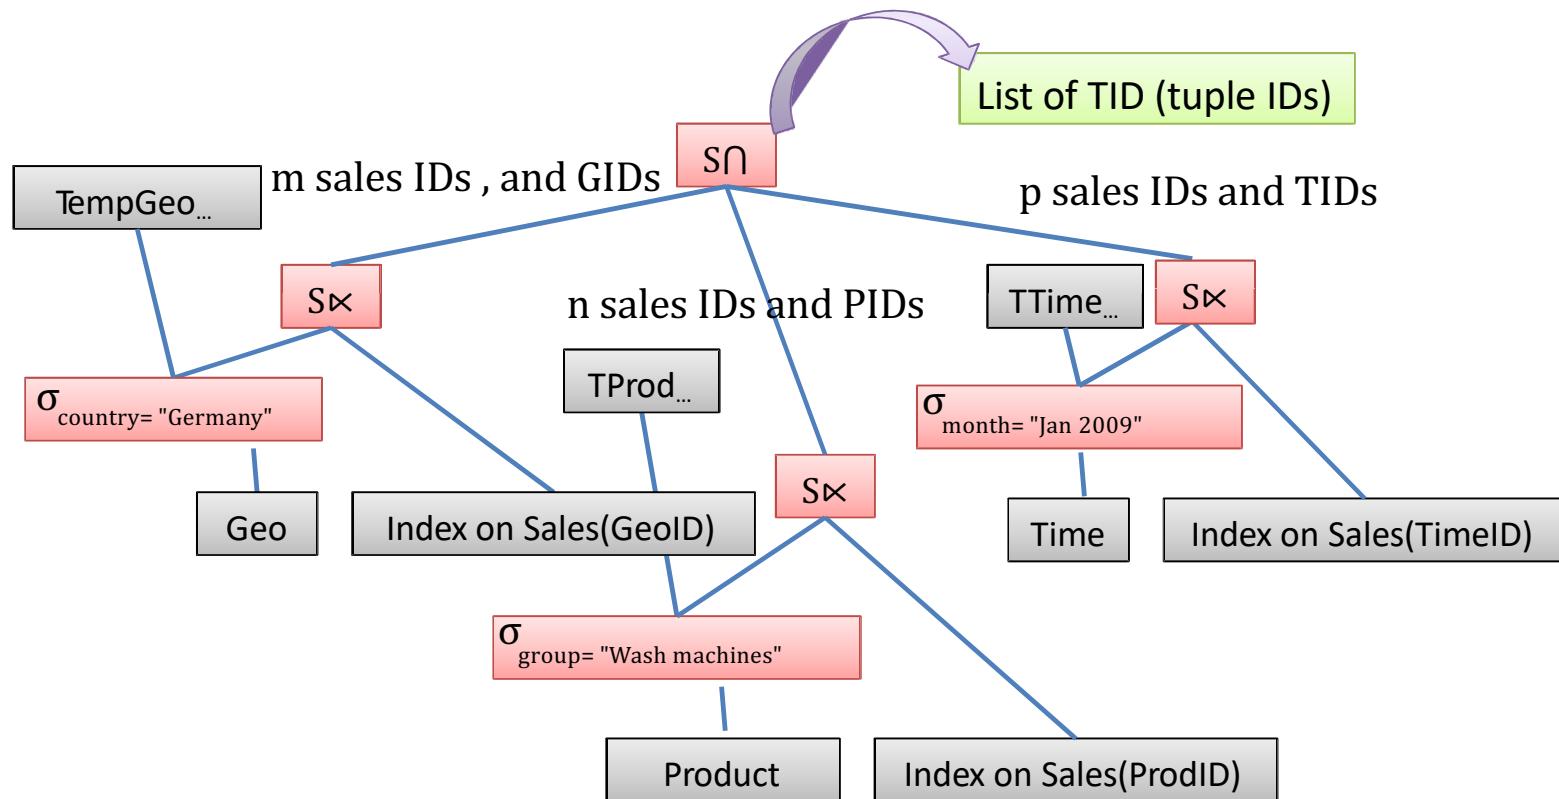


Geoid	ID
1	1
1	2
4	3
...	...

Geoid	ID
1	1
1	2

Star-join Optimization (cont'd.)

- Reduces the fact table to what we need, based on indexes



Materialized Views

- MaterializedViews (MV)
 - Views whose tuples are **stored** in the database are said to be materialized
 - They provides fast access, like a (very high-level) cache
 - Need to maintain the view as the underlying tables change
 - Ideally, we want incremental view maintenance algorithms



Materialized Views (cont'd.)

- How can we use MV in DW?
 - E.g., we have queries requiring us to join the Sales table with another table and aggregate the result
 - `SELECT P.Categ, SUM(S.Qty) FROM Product P, Sales S WHERE P.ProdID=S.ProdID GROUP BY P.Categ`
 - `SELECT G.Store, SUM(S.Qty) FROM Geo G, Sales S WHERE G.Geoid=S.Geoid GROUP BY G.Store`
 -
 - There are more solutions to speed up such queries
 - Pre-compute the two joins involved (product with sales and geo with sales)
 - Pre-compute each query in its entirety
 - Or use an already materialized view

Materialized Views (cont'd.)

- Having the following view materialized
 - CREATE MATERIALIZED VIEW Totalsales (ProdID, GeoID, total) AS SELECT S.ProdID, S.GeoID, SUM(S.Qty) FROM Sales S GROUP BY S.ProdID, S.GeoID
- We can use it in our 2 queries
 - ```
SELECT P.Categ, SUM(T.Total) FROM Product P, Totalsales T WHERE P.ProdID=T.ProdID GROUP BY P.Categ
```
  - ```
SELECT G.Store, SUM(T.Total) FROM Geo G, Totalsales T WHERE G.GeoID=T.GeoID GROUP BY G.Store
```

Materialized Views (cont'd.)

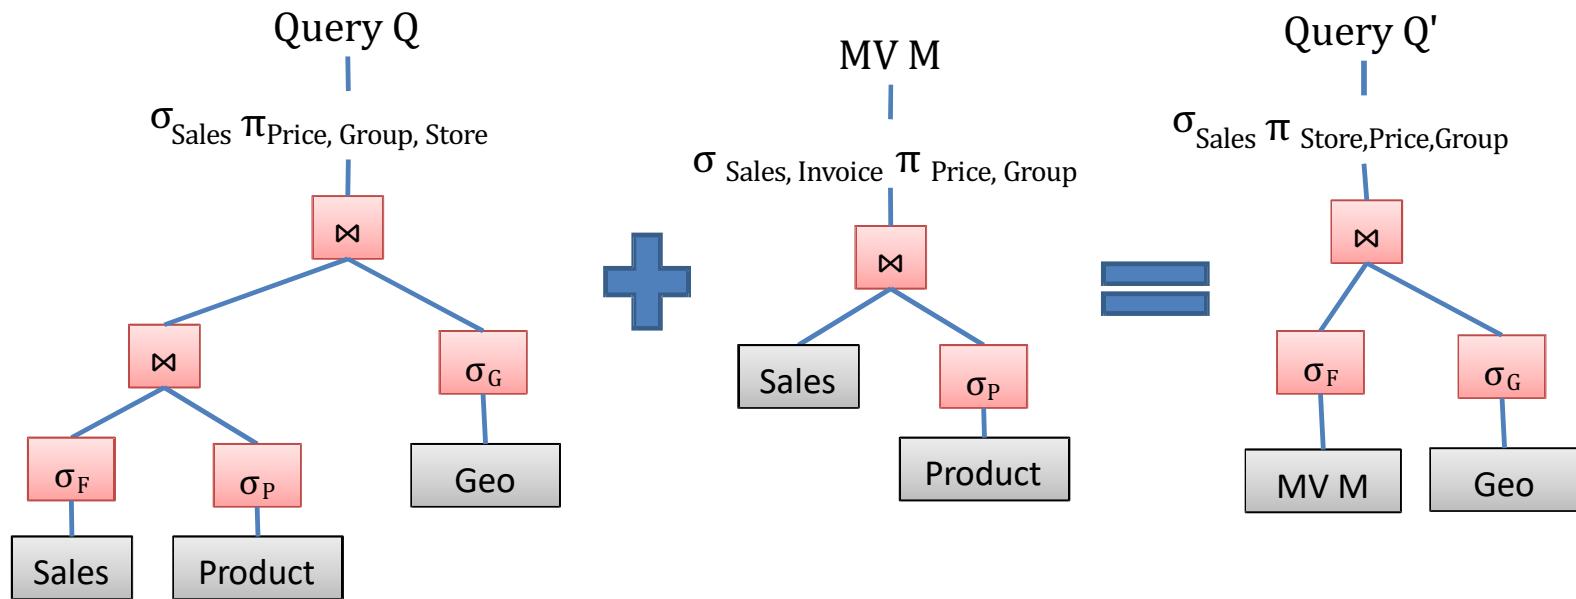
- MV issues
 - **Utilization**
 - What views should we materialize, and what indexes should we build on the pre-computed results?
 - **Choice of materialized views**
 - Given a query and a set of materialized views, can we use the materialized views to answer the query?
 - **Maintenance**
 - How frequently should we refresh materialized views to make them consistent with the underlying tables?
 - And how can we do this incrementally?

Utilization of MV

- Materialized views **utilization** has to be transparent
 - Queries are internally rewritten to use the available MVs by the **query rewriter**
 - The query rewriter performs integration of the MV based on the **query execution graph**

Utilization of MV (cont'd.)

- E.g., materialized views utilization, mono-block query



Integration of MV

- Integration of MV
 - **Valid replacement:** A query Q' represents a valid replacement of query Q by utilizing the materialized view M , if Q and Q' always deliver the same result set
 - For general relational queries, the problem of finding a valid replacement is **NP-complete**

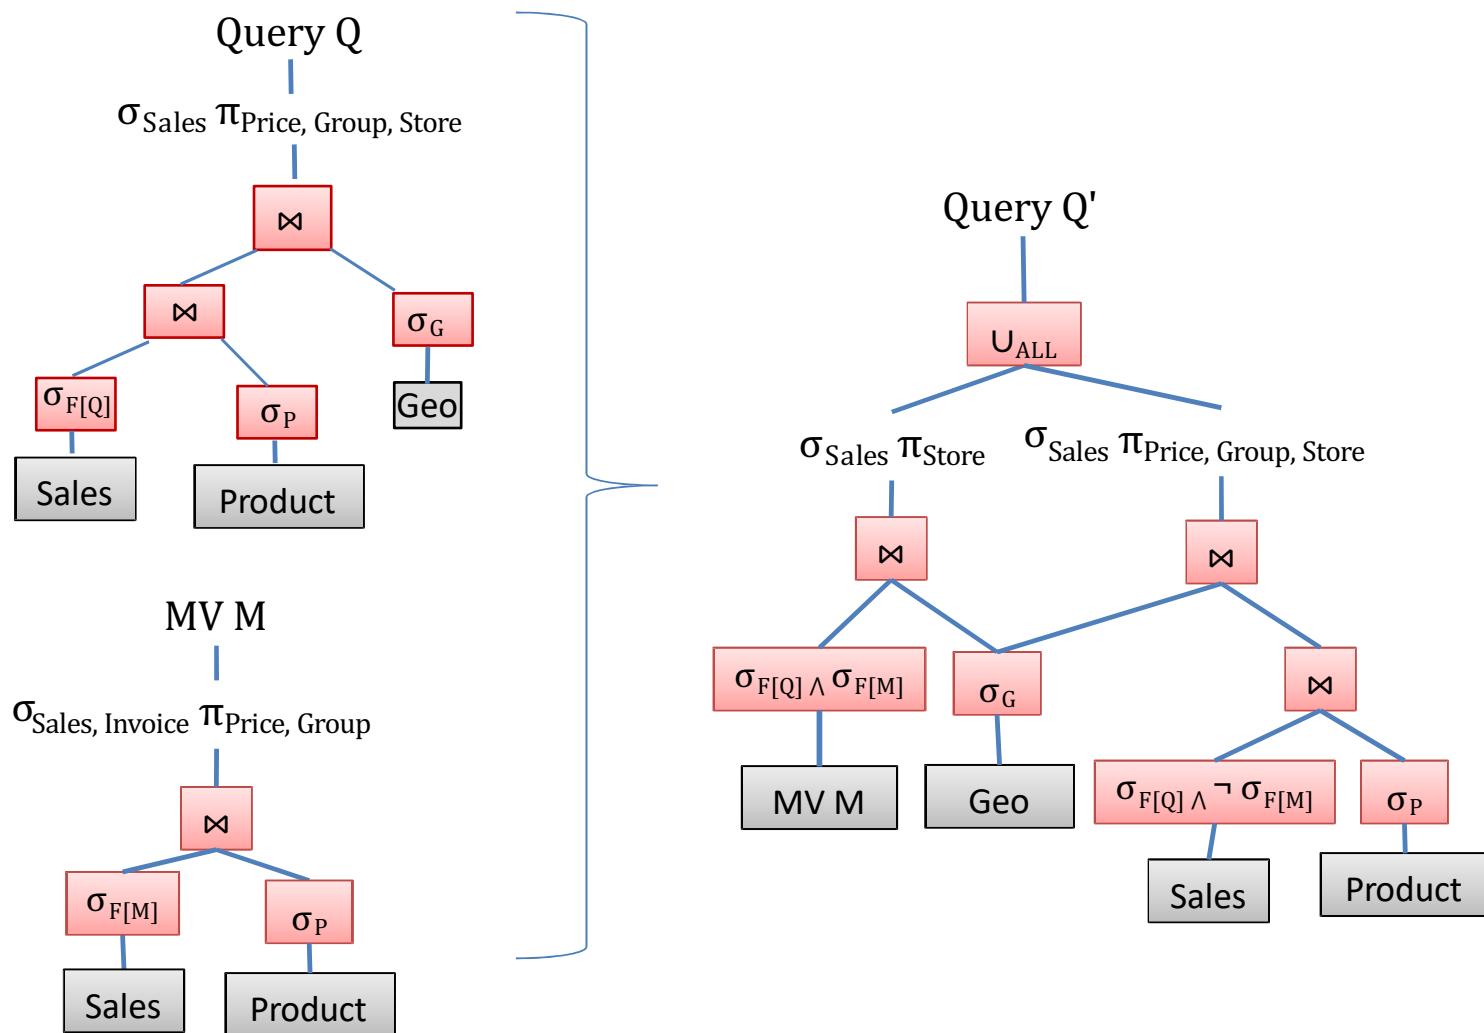
Integration of MV (cont'd.)

- In order to be able to integrate MV M in Q and obtain Q', the following **conditions** need to be respected
 - The selection condition in M cannot be more restrictive than the one in Q
 - The projection from Q has to be a subset of the projection from M
 - It has to be possible to derive the aggregation functions of $\pi(Q)$ from $\pi(M)$
 - Additional selection conditions in Q have to be possible also on M

Integration of MV (cont'd.)

- How do we use MV even when there is no perfect match?
- If the selection in M is more restrictive than the selection in Q
 - Split Q in Q_a and Q_b such that $\sigma(Q_a) = (\sigma(Q) \wedge \sigma(M))$ and $\sigma(Q_b) = (\sigma(Q) \wedge \neg \sigma(M))$
 - Multiblock queries

Integration of MV (cont'd.)



Summary

Summary

- Partitioning: Horizontal or Vertical
 - Records used together are grouped together
 - However: slow retrieval across partitions
 - Mini-Dimensions
- Joins: for DW it is sometimes better to perform cross product on dimensions first
- Materialized Views: we can't materialize everything
 - Static or Dynamic choice of what to materialize
 - The benefit cost function is decisive

Next Lecture

- Queries!
 - OLAP queries
 - SQL for DW
 - MDX

