
Lecture 11

Real-time Data Warehousing


Summary – last week

Summary

- How to build a DW
 - The DW Project: usual tasks, hardware, software, timeline (phases)
 - Data Extract/Transform/Load (ETL):
 - Data storage structures, extraction strategies (e.g., scraping, sniffing)
 - Transformation: data quality, integration
 - Loading: issues, and strategies, (bulk loading for fact data is a must)
 - Metadata:
 - Describes the contents of a DW, comprises all the intermediate products of ETL,
 - Helps for understanding how to use the DW

This week

- Real-time Data Warehousing
 - Real-time Data Processing Challenges
 - Real-time Join Algorithms



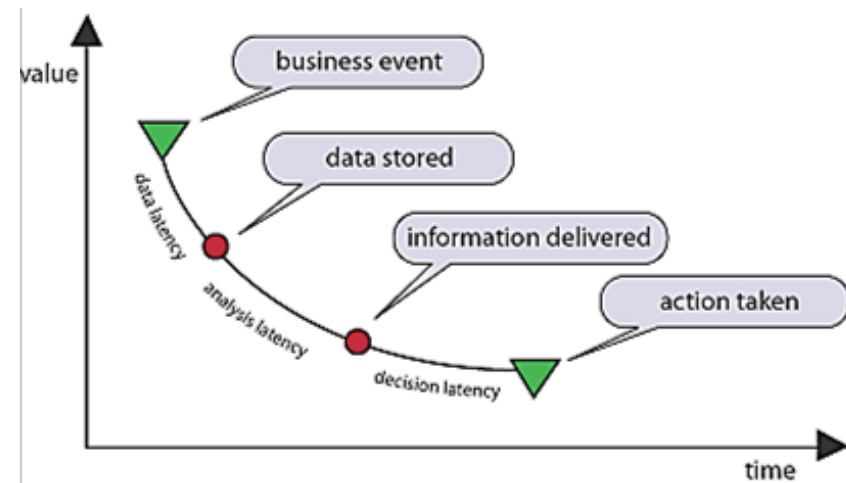
Real-time DW – Why?

- **Customer business scenario:** a utility company owns plants generating energy
- Existing DW supports planning by recommending:
 - The production capacity
 - The reserve capacity
 - When to buy supplemental energy, as needed



Real-time DW – Why?

- Each day is pre-planned on **historical behavior**
 - Peak demand periods are somewhat predictable
- Good planning is important because:
 - Expensive to have unused capacity!
 - Cheaper to buy energy ahead!
- Planning on last week's average is not enough



Real-time DW – Why?

- Getting more **in-time** accuracy enhances operational business
 - Compare **today's** plant output and customer consumption volumes to yesterday's or last week's average
 - Know when to purchase additional options or supplies
- Customer Target: have the actual data from the operational environment available for analytics within a **5 minute lag**
- **Real-time \neq fast**
 - Real time DW has the capability to **enforce** time constraints


Real-time DW – Why?

- The most difficult part for complying with the 5 minutes time constraint is the **ETL process**
 - ETL tools usually operate in batch mode on a certain schedule nightly, weekly or monthly
 - ETL typically involves **downtime** for the DW during the loading step (usually happens over night)
 - Data is extracted into flat files in the staging area **outside the DBMS**, where it is not available for querying
 - ETL may take **hours** to finish
- The ETL process needs to be re-designed to meet real-time constraints

Real-time Data Warehousing

- Real-Time Data Warehouse (RTDW)


Individual changes occurring on the source systems are immediately forwarded to data warehouse in a best effort strategy.



- RTDW needs Real-Time ETL (Extract, Transform, Load)

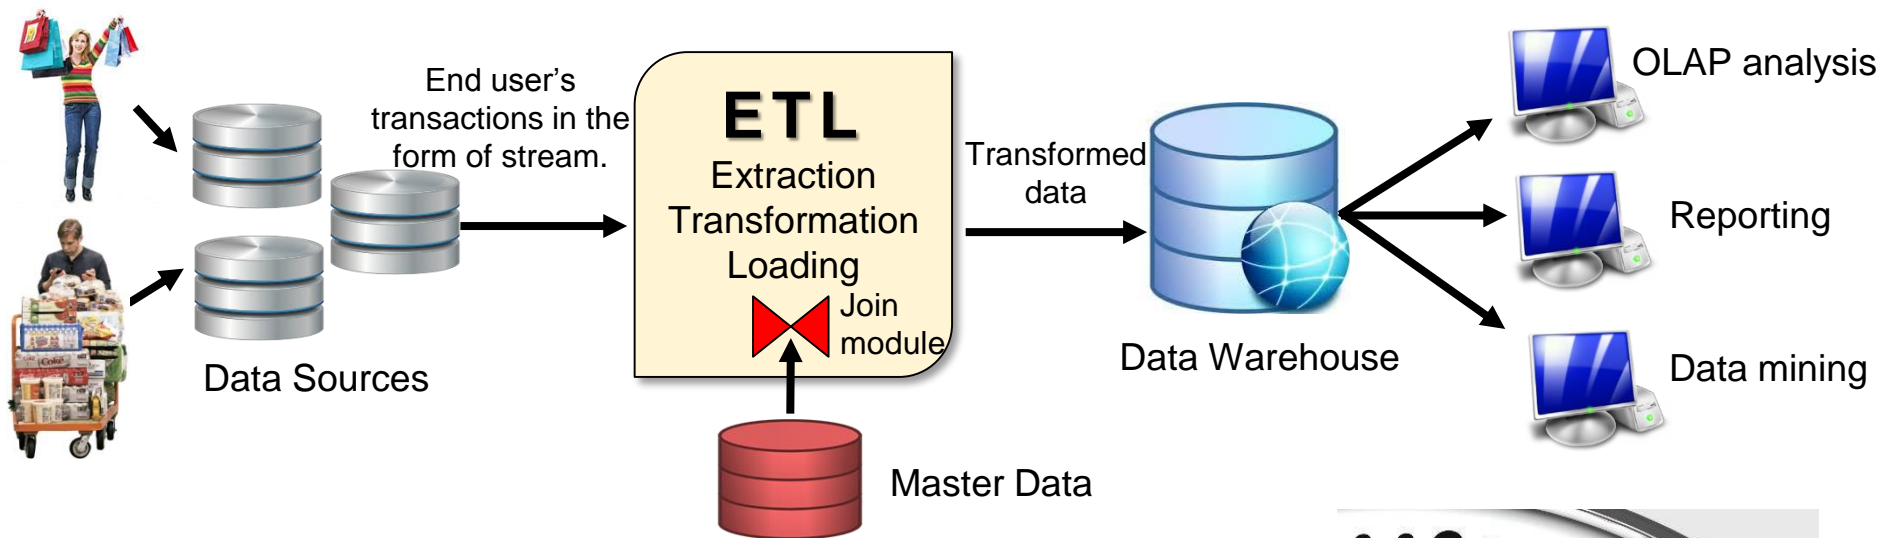
- ETL is a data integration layer between multi-data sources and data warehouse.

- Perform Extraction, Transformation and Loading tasks on Near-Real-Time basis



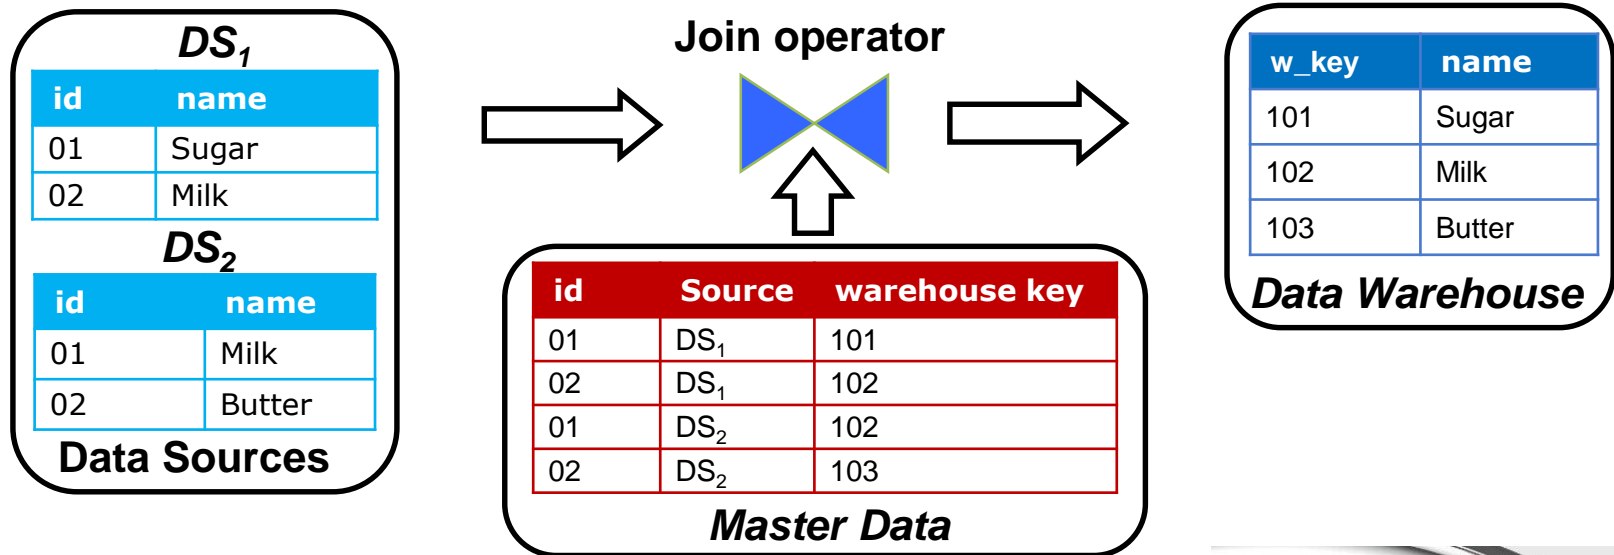
Real-time ETL

Real-time transformation is an important phase in an ETL layer where incoming source updates are transformed into warehouse format in an online fashion.





Real-time ETL (cont'd.)

Transformation examples: Key replacement, enrichment of data



What is Stream-based Join?

- Stream-based join is an operation to combine the information coming from two or more data sources.
- Data sources can be in the form of streams or persistent data.



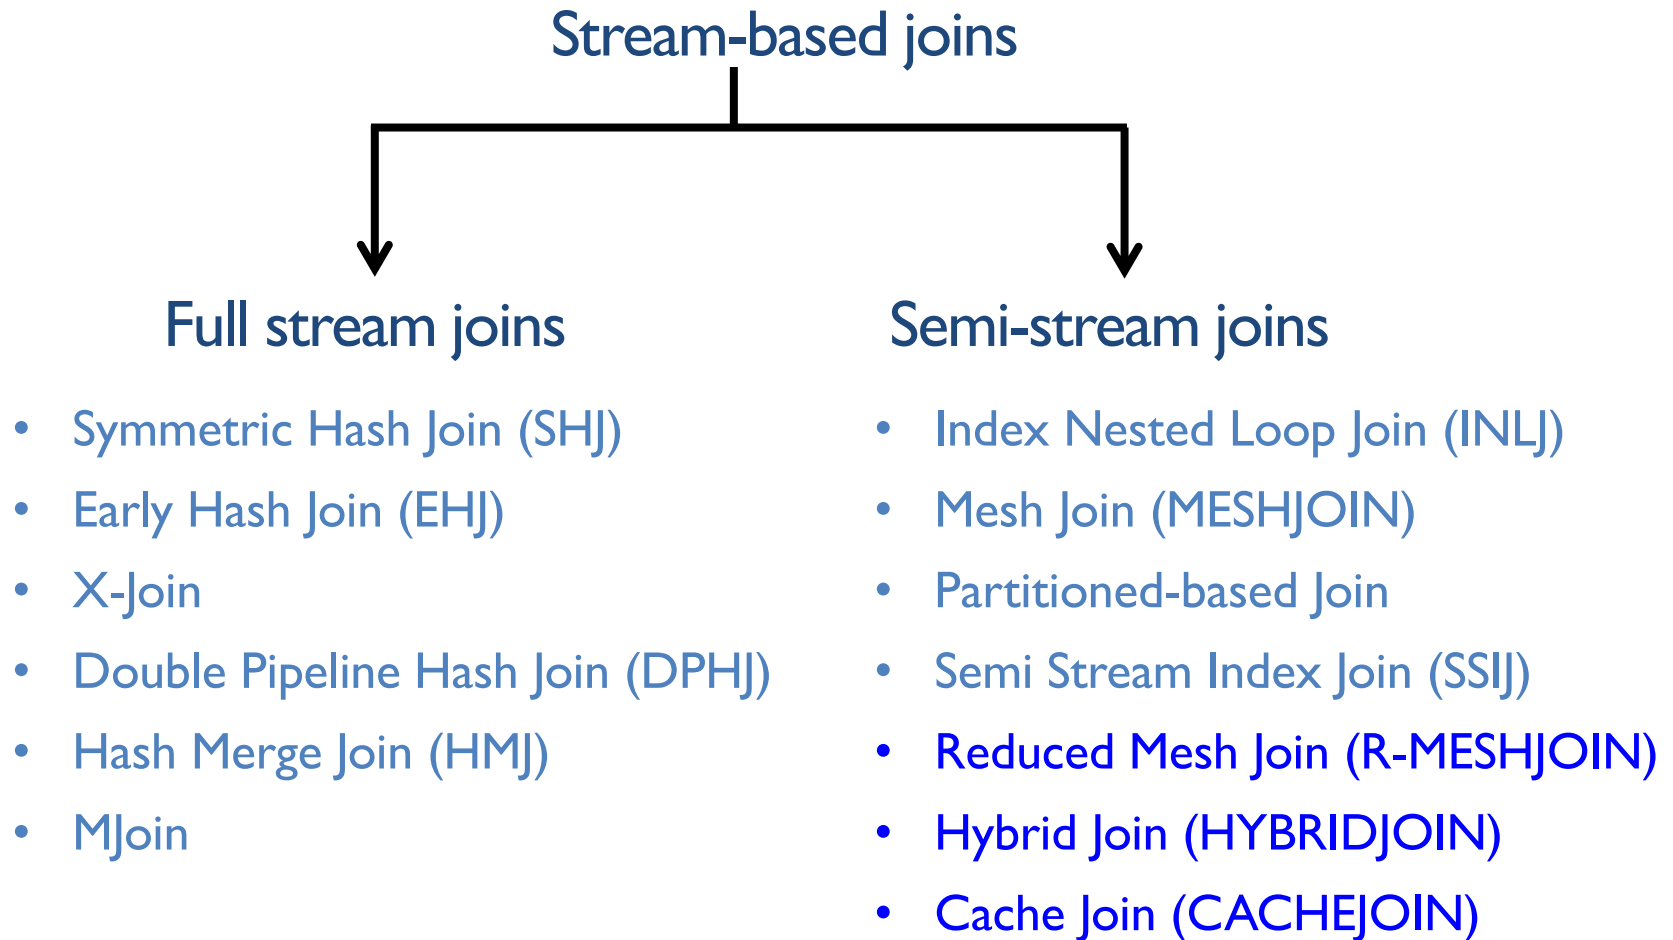
- Applications of stream-based join.
 - Enrichment of stream data with master data.
 - Key replacement in data warehouse.
 - Identification of duplicate records.
 - Merging of two or more data streams.




Research Challenges

- Challenge 1
 - Both inputs of the join operator have different arrival rates.
 - The stream input is fast, high volume and has an intermittent nature.
 - Master data input is comparatively slow due to the disk I/O cost.
 - It creates a bottleneck in join processing. The challenge here is to eliminate this bottleneck.
- Challenge 2:
 - Stream data is non-uniform therefore, an efficient approach is required to retrieve master data.

Existing Approaches



Index Nested Loop Join (INLJ)



Issues in INLJ

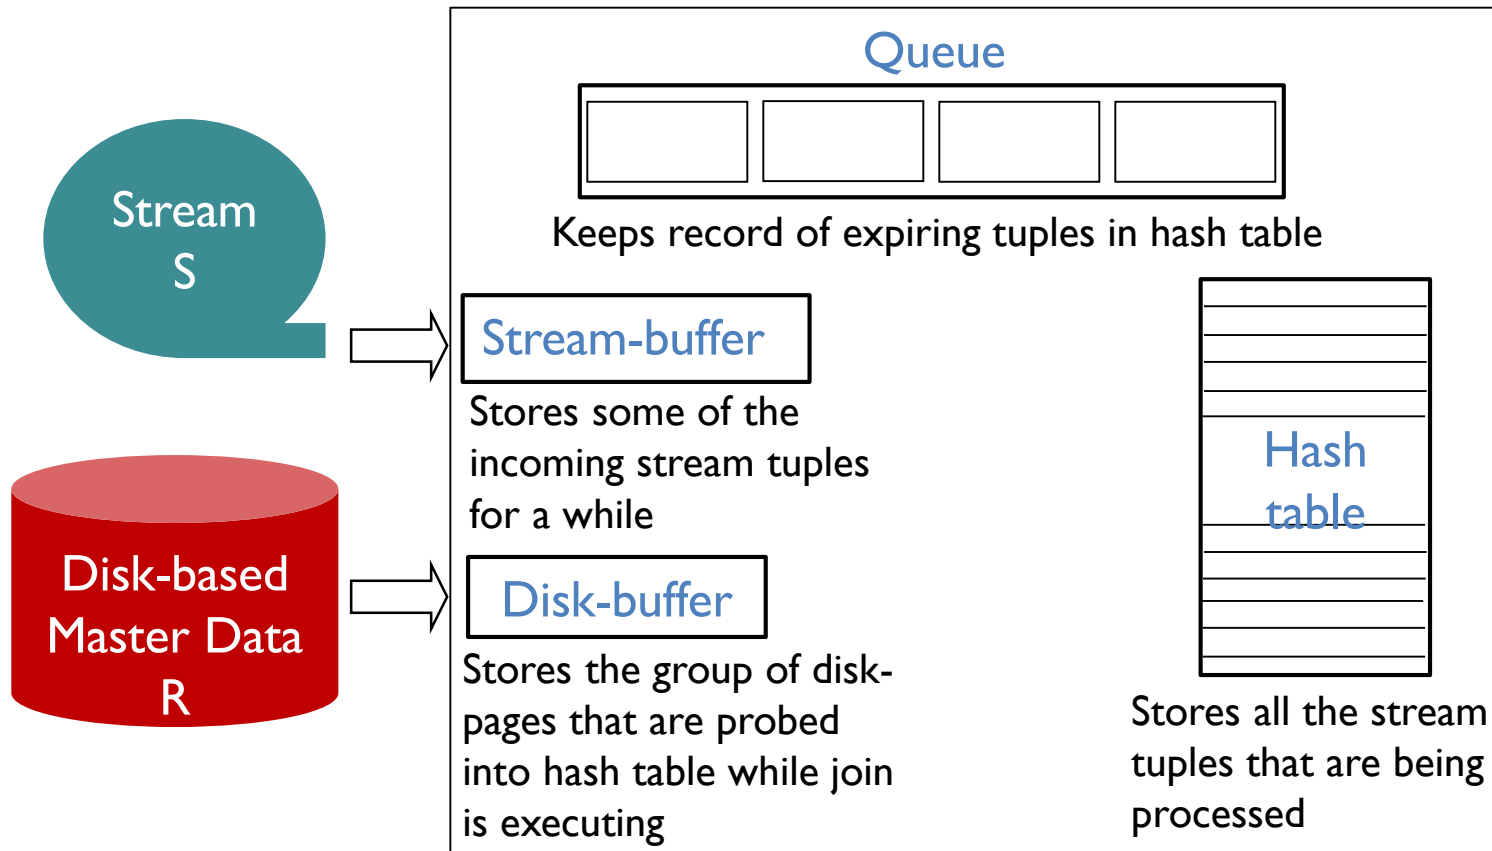
- INLJ processes only one stream tuple per iteration therefore solution is not practical for fast and huge volume stream data.
- INLJ does not amortize expensive disk reading cost on stream data.
- INLJ does not take into account some common characteristics of stream data e.g. skew in stream data.

MESHJOIN (Mesh Join)

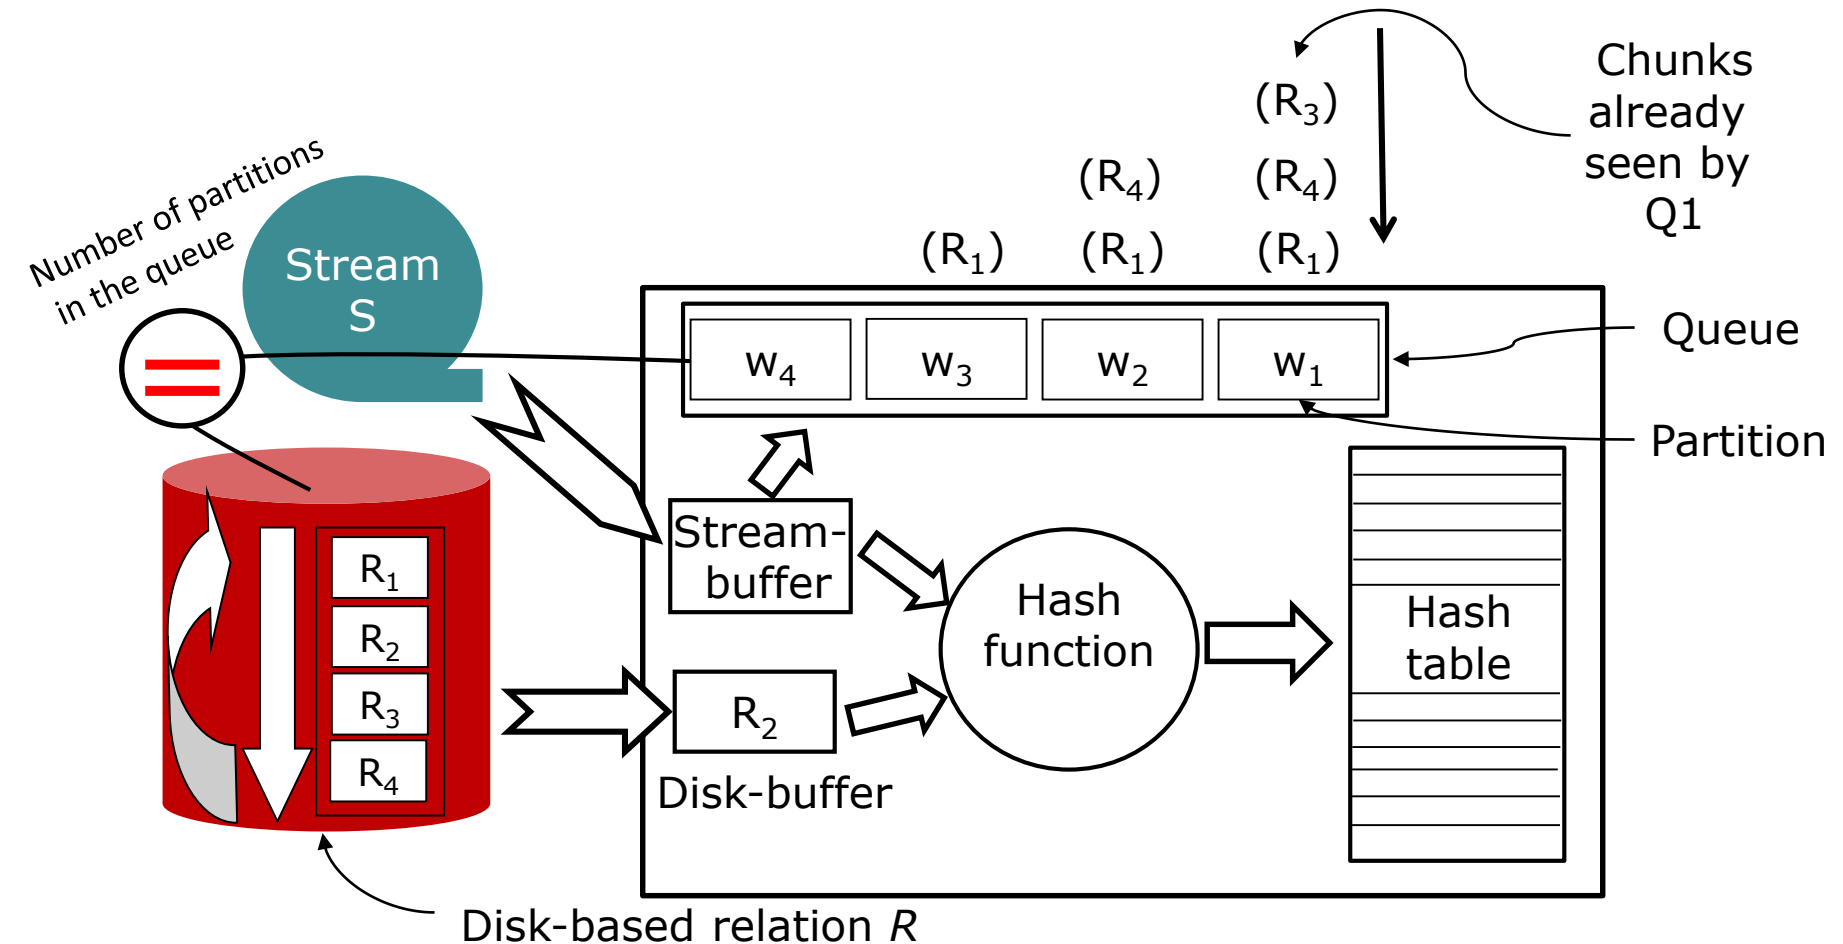
Features

- MESHJOIN (Mesh Join) has been proposed for processing stream data with master data.
- Designed for joining a stream S with disk-based relation R .
- Uses limited memory budget.
- Does not need an index.
- Works for any equijoin.

MESHJOIN Components




MESHJOIN Operation



- Size of stream-buffer = w tuples
- Size of disk-buffer = b pages
- Iterations required to bring all of R into memory = k (in this example $k=4$)

Problem in MESHJOIN

- Problem 1
 - Due to unnecessary dependency, memory distribution among the join components is not optimal.
- Problem 2
 - The performance of the algorithm is inversely proportional to the size of master data.
- Problem 3
 - Can not deal with intermittency in stream data.
- Problem 4
 - Typical characteristics of stream data such as non-uniform data are not considered.

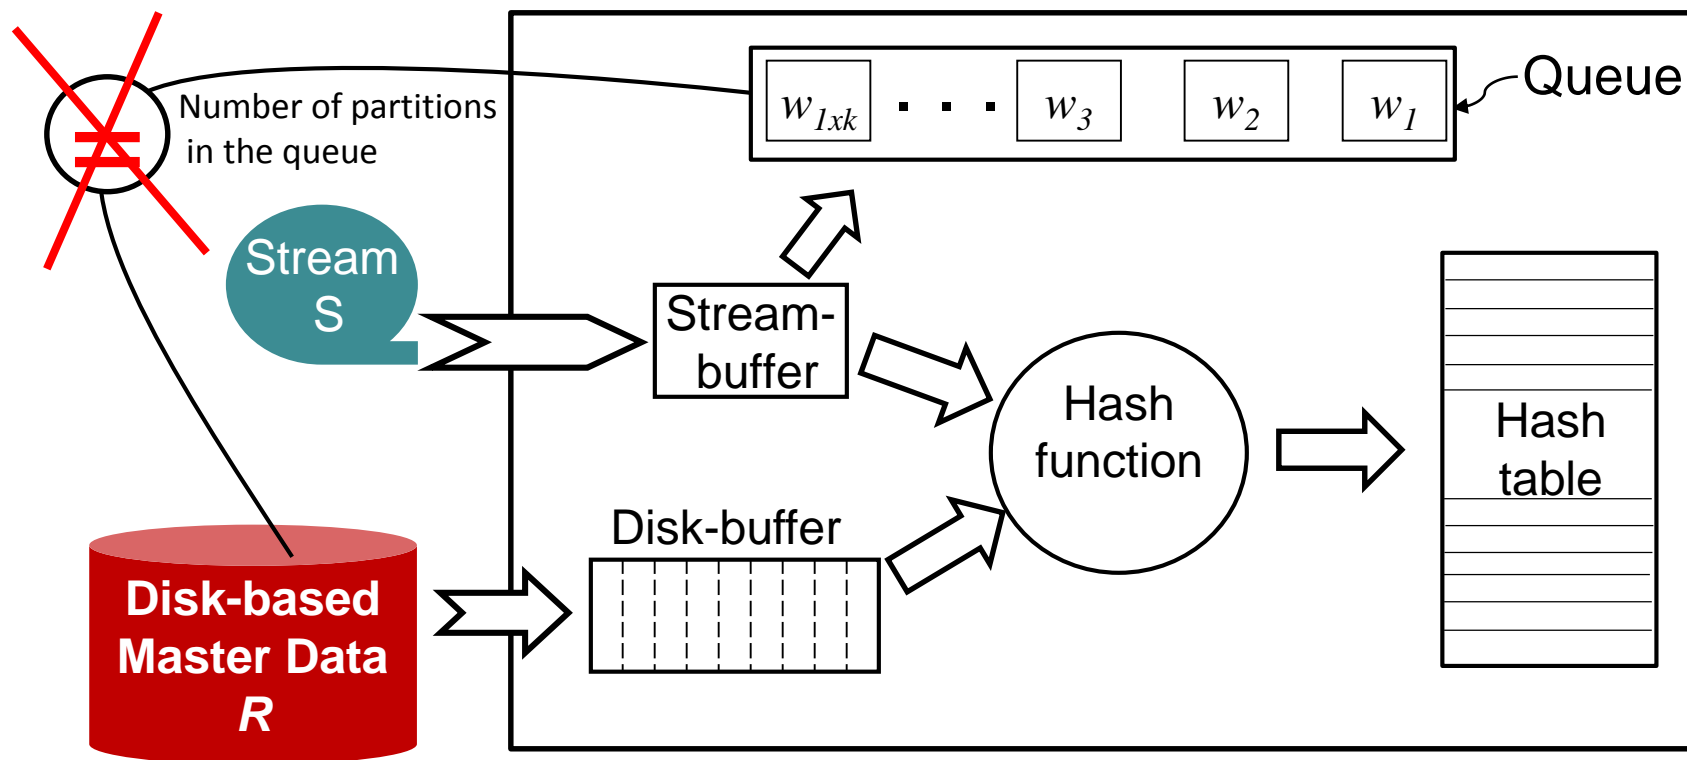


Solutions

We propose following novel algorithms to solve the highlighted problems.


- **Reduced Mesh Join (R-MESHJOIN):-** clarifies the dependency among the components more appropriately.
- **Hybrid Join (HYBRIDJOIN):-** introduces an index-based strategy to access the master data.
- **Cache Join (CACHJOIN):-** considers non-uniform characteristic in stream data.

R-MESHJOIN (Reduced Mesh Join)



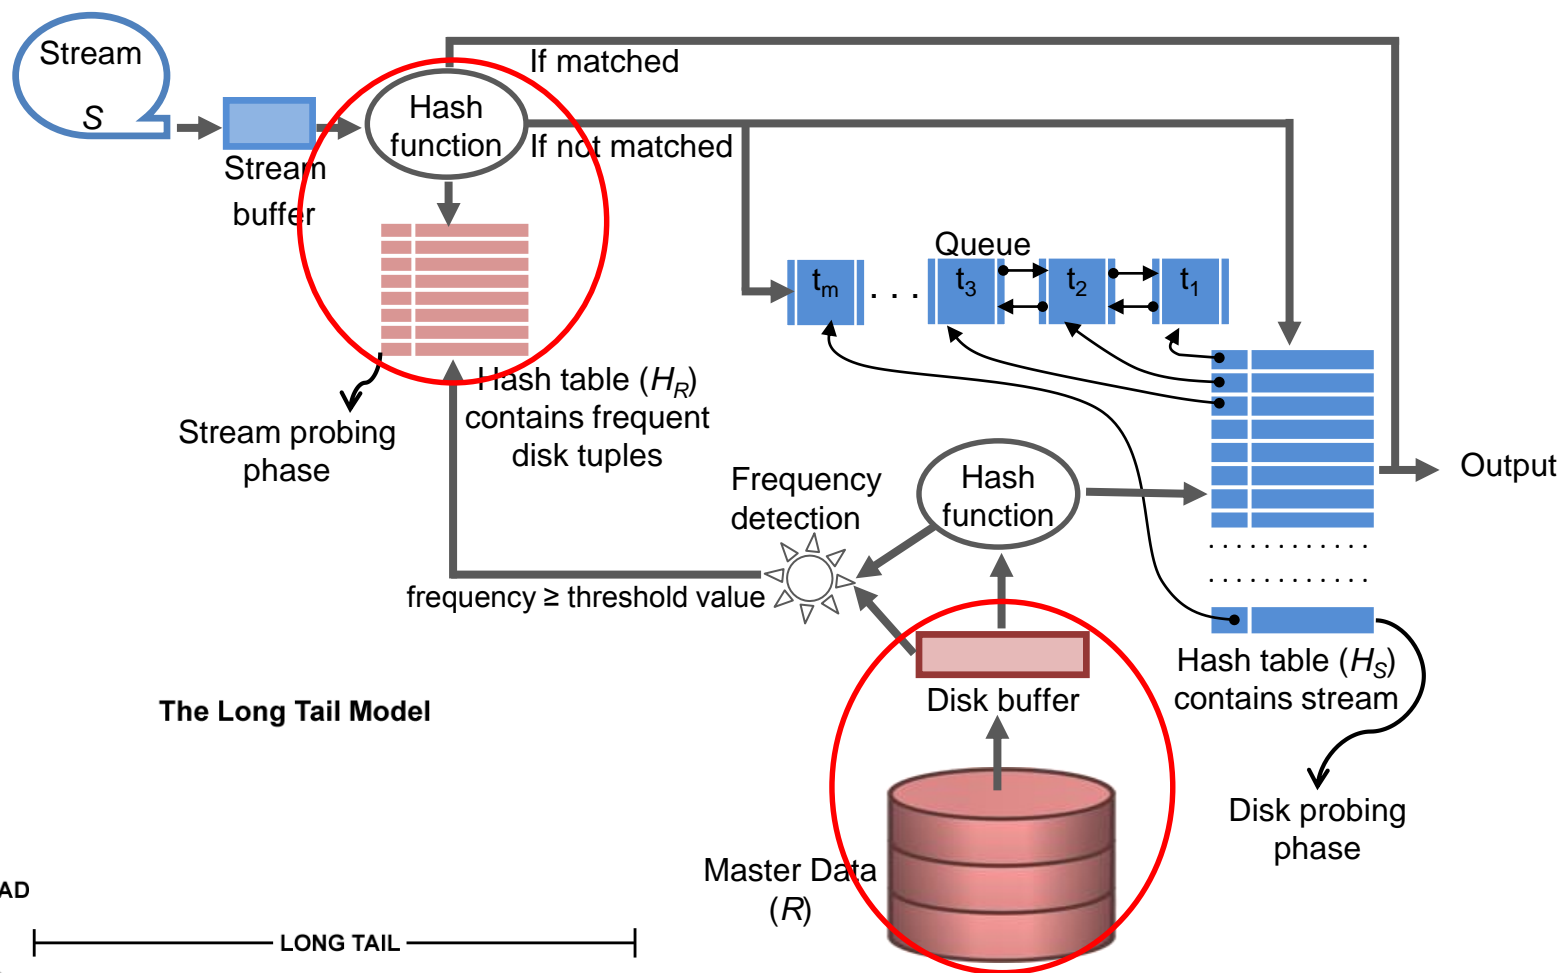
- Number of logical partitions in disk-buffer = l
- Size of each logical partition (pages) = b_p
- Size of disk-buffer (pages) = b
- Iteration required to load entire disk-based relation into memory = k

HYBRIDJOIN (Hybrid Join)



1. Assume that R is sorted in the order of access frequency.
2. Index-based strategy is used to access R.

CACHEJOIN (Cache Join)



Experimental Setup

- Hardware specifications
 - Pentium-core-i5, 8G main memory, 500G HDD
- Synthetic dataset
 - Size of master data R, 100 million tuples (~11.18GB)
 - Size of each disk tuple, 120 bytes
 - Size of each stream tuple, 20 bytes
 - Size of each pointer in queue, 12 bytes
 - Based on Zipf's Law with skew of 0 to 1

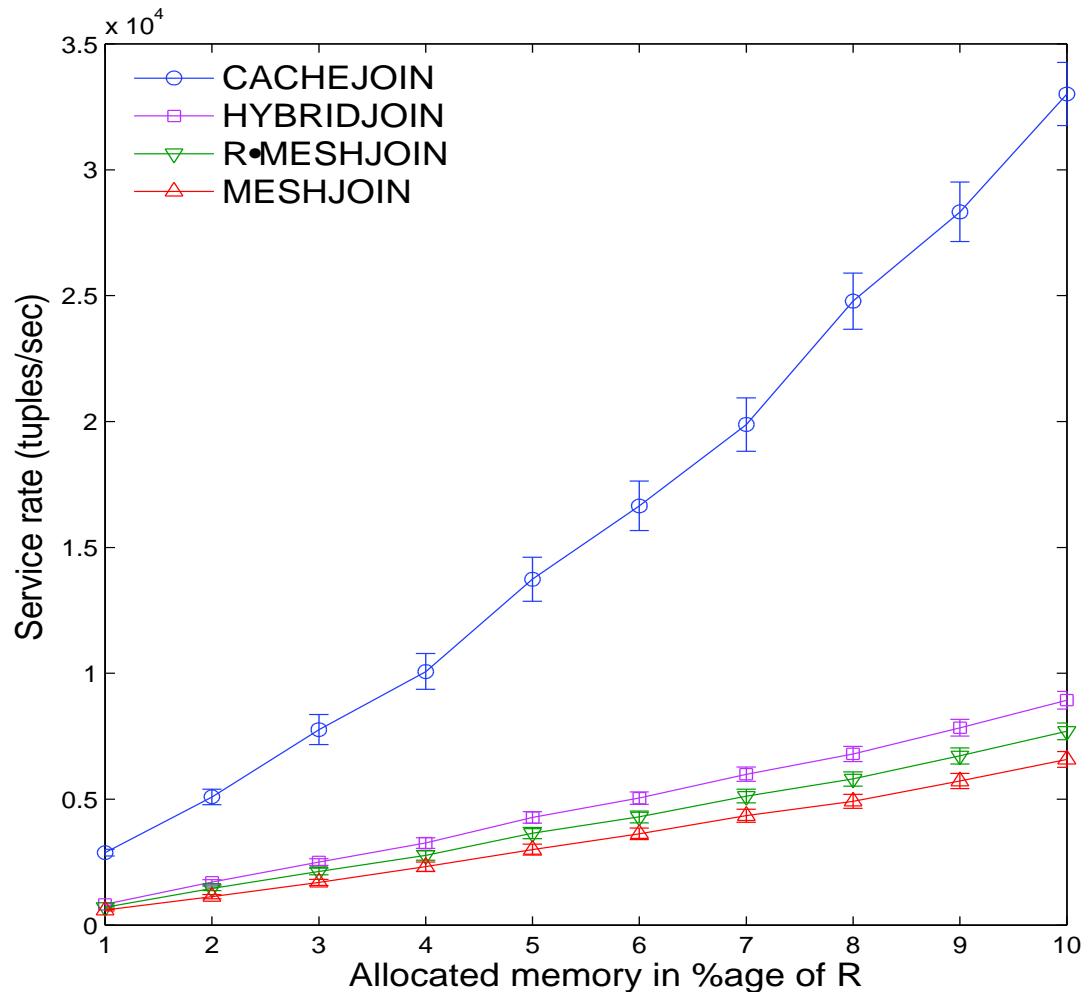


Experimental Setup (cont'd.)

- TPC-H dataset
 - Create the datasets using a scale factor of 100
 - Uses table Customer as our master data table with each tuple of 223 bytes
 - Uses table Order as our stream data table with each tuple of 138 bytes
- Real dataset
 - Size of master data, 20 million tuples
 - Size of each master as well as stream tuple is 128 bytes
 - Source url: <http://cdiac.ornl.gov/ftp/ndp026b/>
- Evaluation metrics
 - We calculate confidence interval by considering 95% accuracy rate.

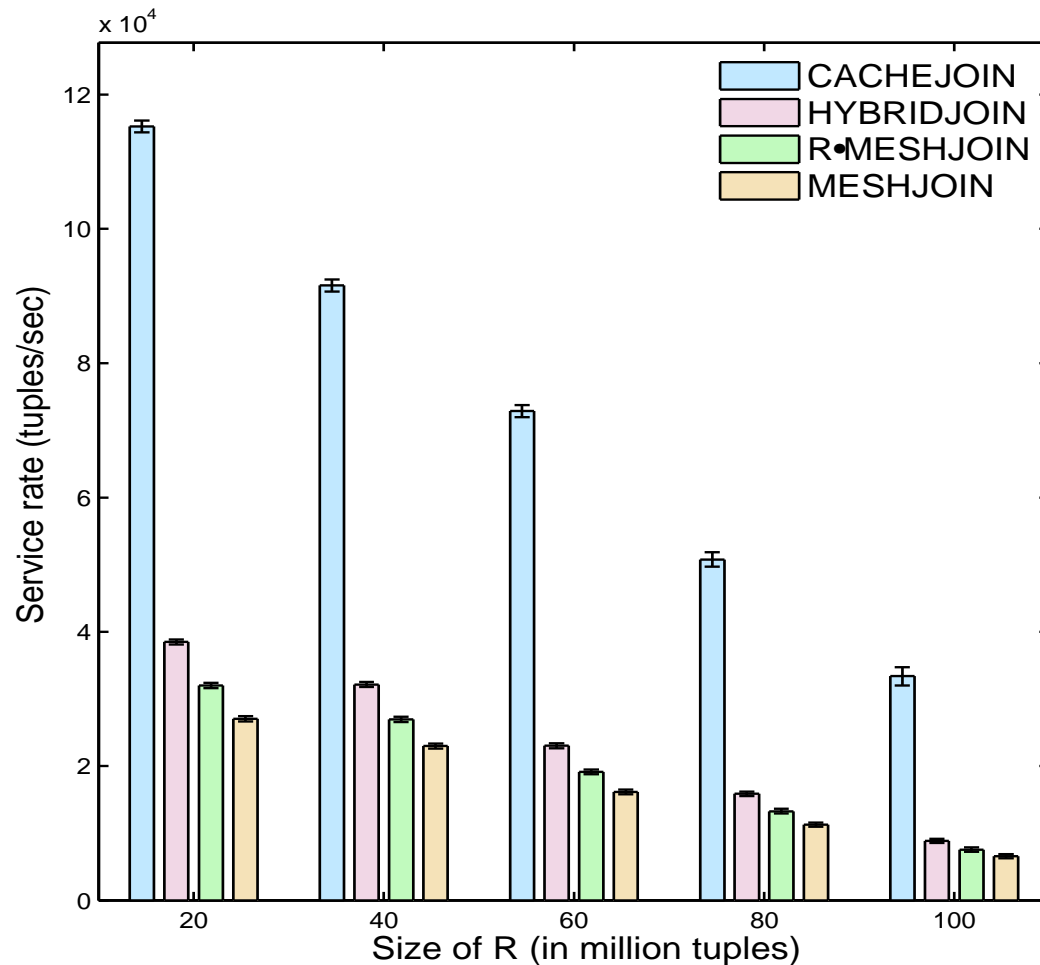
Performance Comparisons

Performance comparisons with 95% confidence interval while $R=100$ million tuples and **M** varies in percentage of R .




Performance Comparisons (cont'd.)

Performance comparisons with 95% confidence interval while $M \sim 1.2G$ and **R varies**.



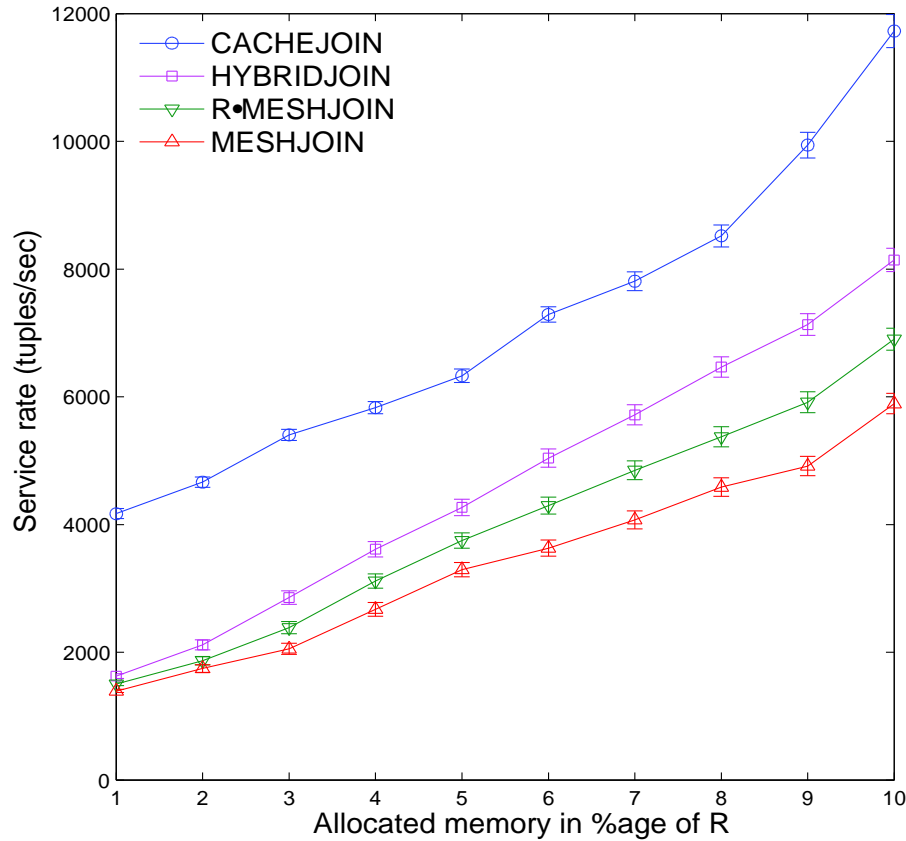
Performance Comparisons (cont'd.)

Performance comparisons with 95% confidence interval while $M \sim 1.2G$, $R = 100$ million tuples and **skew varies**.




Performance Comparisons (cont'd.)

TPC-H datasets

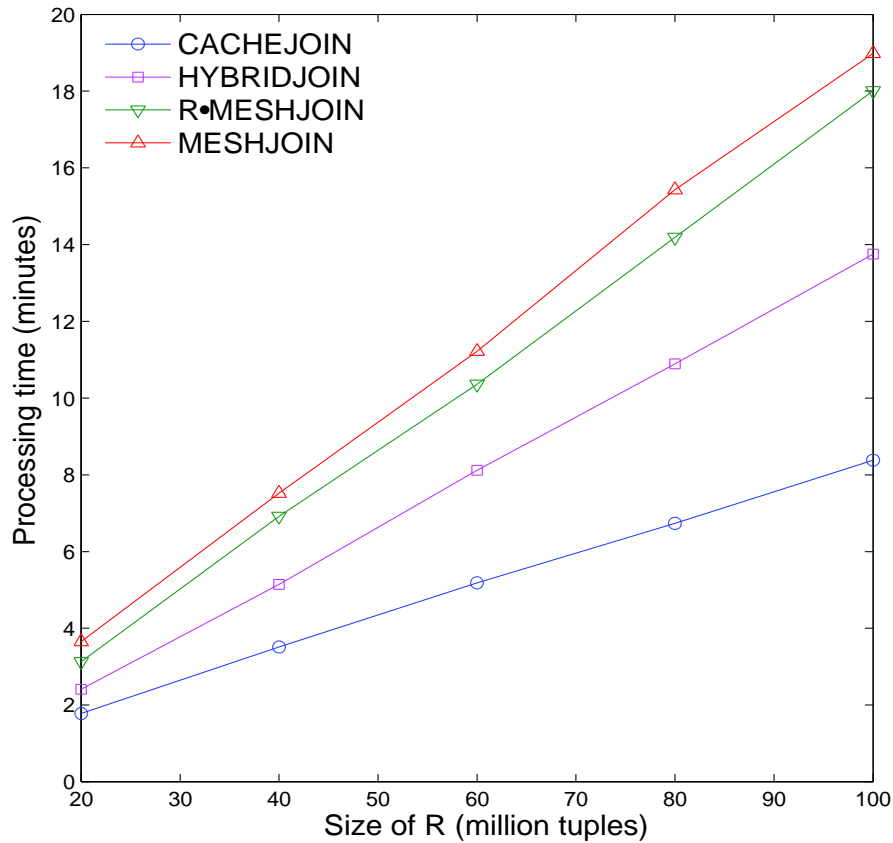


Real datasets

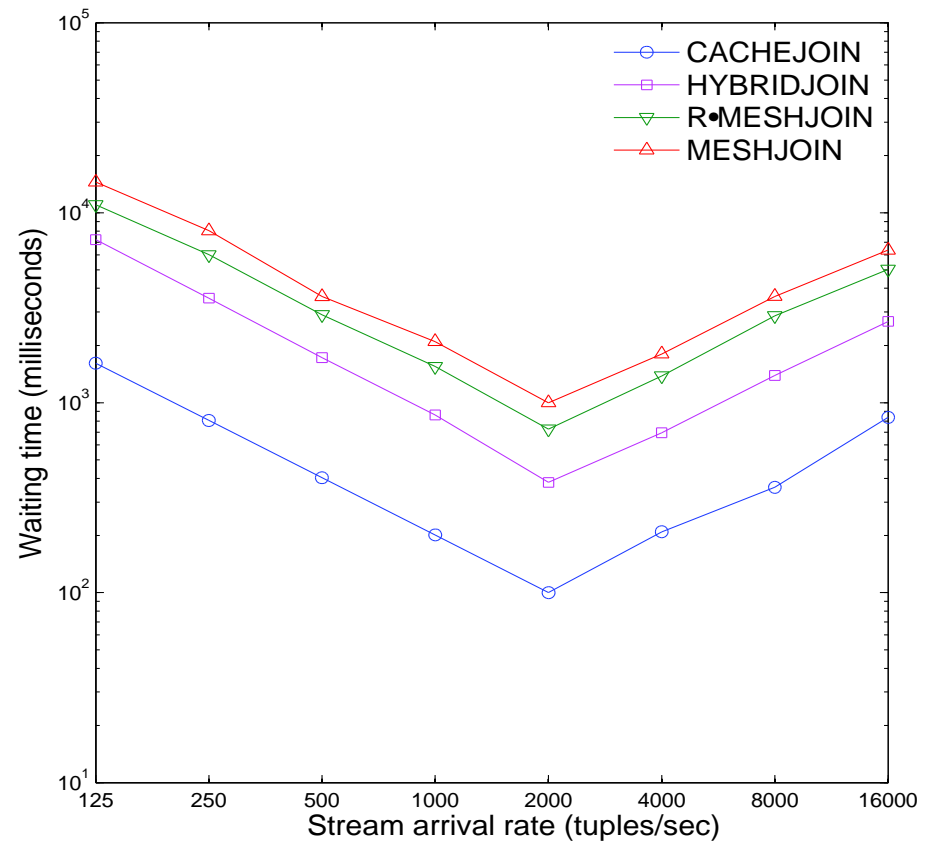


Performance Comparisons (cont'd.)

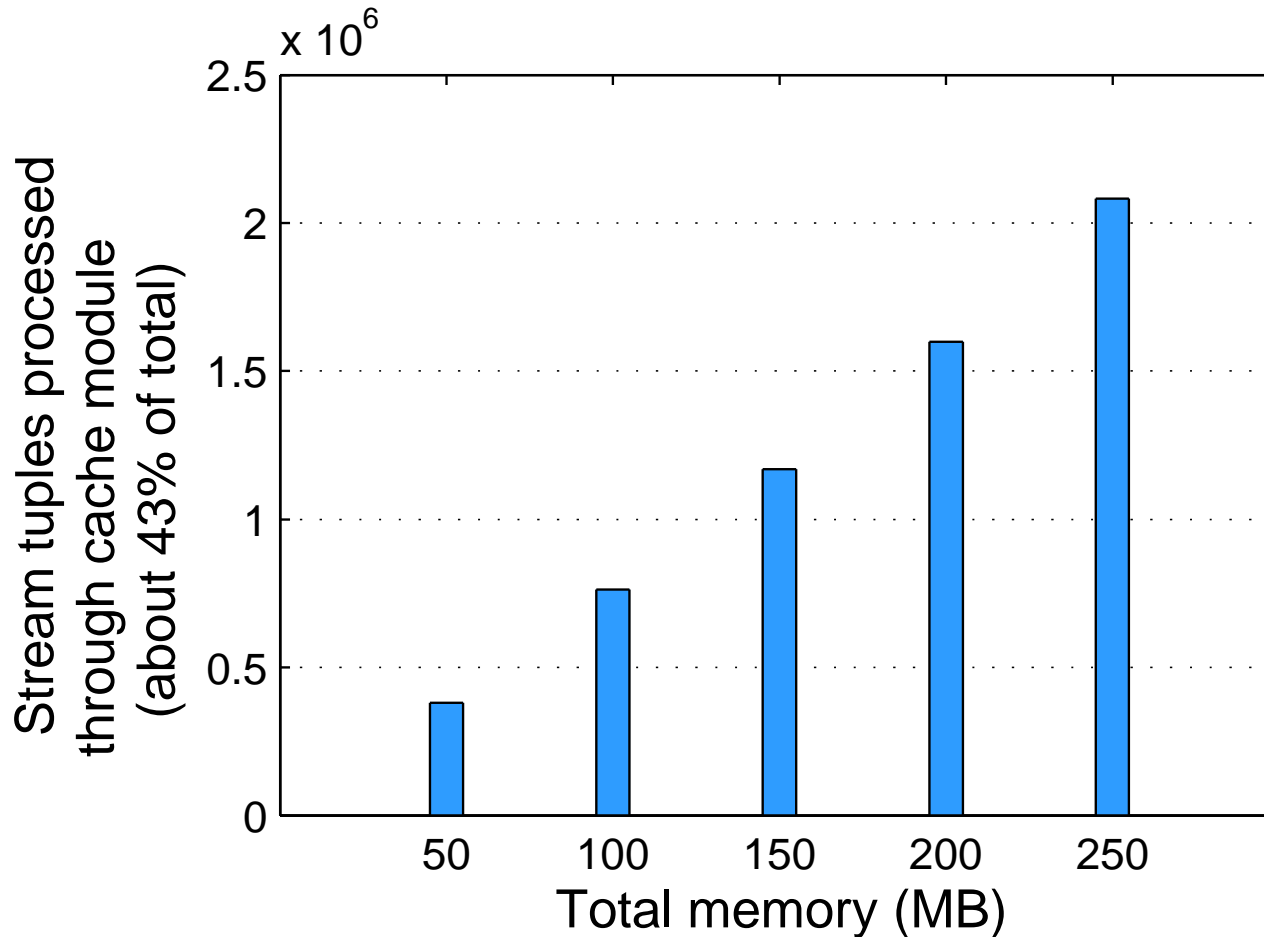
Processing time



Waiting time



Role of Cache Module in Performance



Total number of stream tuples processed through cache module in 4000 iterations.

Conclusions

- Stream processing has become a novel field in the area of data management due to its infinite characteristics.
- Stream-based join operators perform a key role in processing of stream data.
- A number of algorithms were designed to process semi-stream data however, they suffer with some limitations.
- We addressed these limitations in our research by presenting three novel algorithms.
- Our experimental evaluation proved the contribution of each algorithm in terms of performance.

About Me

- Publications
 - Published one book, six peer-reviewed journal articles, about thirty core-ranked conference and workshop papers, and three book chapters
- Professional Activities.
 - Keynote speaker, ICDIM'13
 - General track chair, ICDIM'13
 - Editorial member in IJES & JCI
 - PC member in ACSE, ADC, DOLAP, and DASFAA
- Student's Supervision
 - Supervising three Master and two PhD students



My Research

- Data stream processing in real-time
Databases and Data Warehouses → **Research in AUT**
Prof. Albert Yeap, Dr. Russel Pears
Shoba Tegginmath
- Stream processing in Cloud
Computing } **Research with
collaboration of
AUT & UoA**
Prof. Gillian Dobbie
Dr. Gerald Weber
Dr. Christof Lutteroth
- Big Data Management
- Natural Language Processing and
Object Modelling → **Research with collaboration of
AUT & Birmingham University, UK**
Dr. Behzad Bordbar, Dr. Imran S. Bajwa



Albert Yeap



Gillian Dobbie



Gerald Weber



Christof Lutteroth



Behzad Bordbar



Imran S. Bajwa



Shoba Tegginmath

Open Research Issues

- Data Stream Processing
 - Consider many-to-many equijoins
 - Dealing with non-equijoins
 - Parallelization of semi-stream joins
 - Semi-stream joins in cloud computing
 - Semi-stream joins on mobile and embedded platforms
- Other Research Areas
 - Big Data Management and Knowledge Engineering
 - Data Mining

Summary

Summary

- Real-time DW
 - Real-time ETL
 - Joins to process stream data
 - MESHJOIN
 - Problems in MESHJOIN
 - R-MESHJOIN
 - HYBRIDJOIN
 - CACHEJOIN
 - Performance comparisons
 - My research

Next Lecture

- Big Data
 - Real-time Data Processing Challenges
 - Real-time Join Algorithms

